

Lecture #10

Containers

Spring 2024

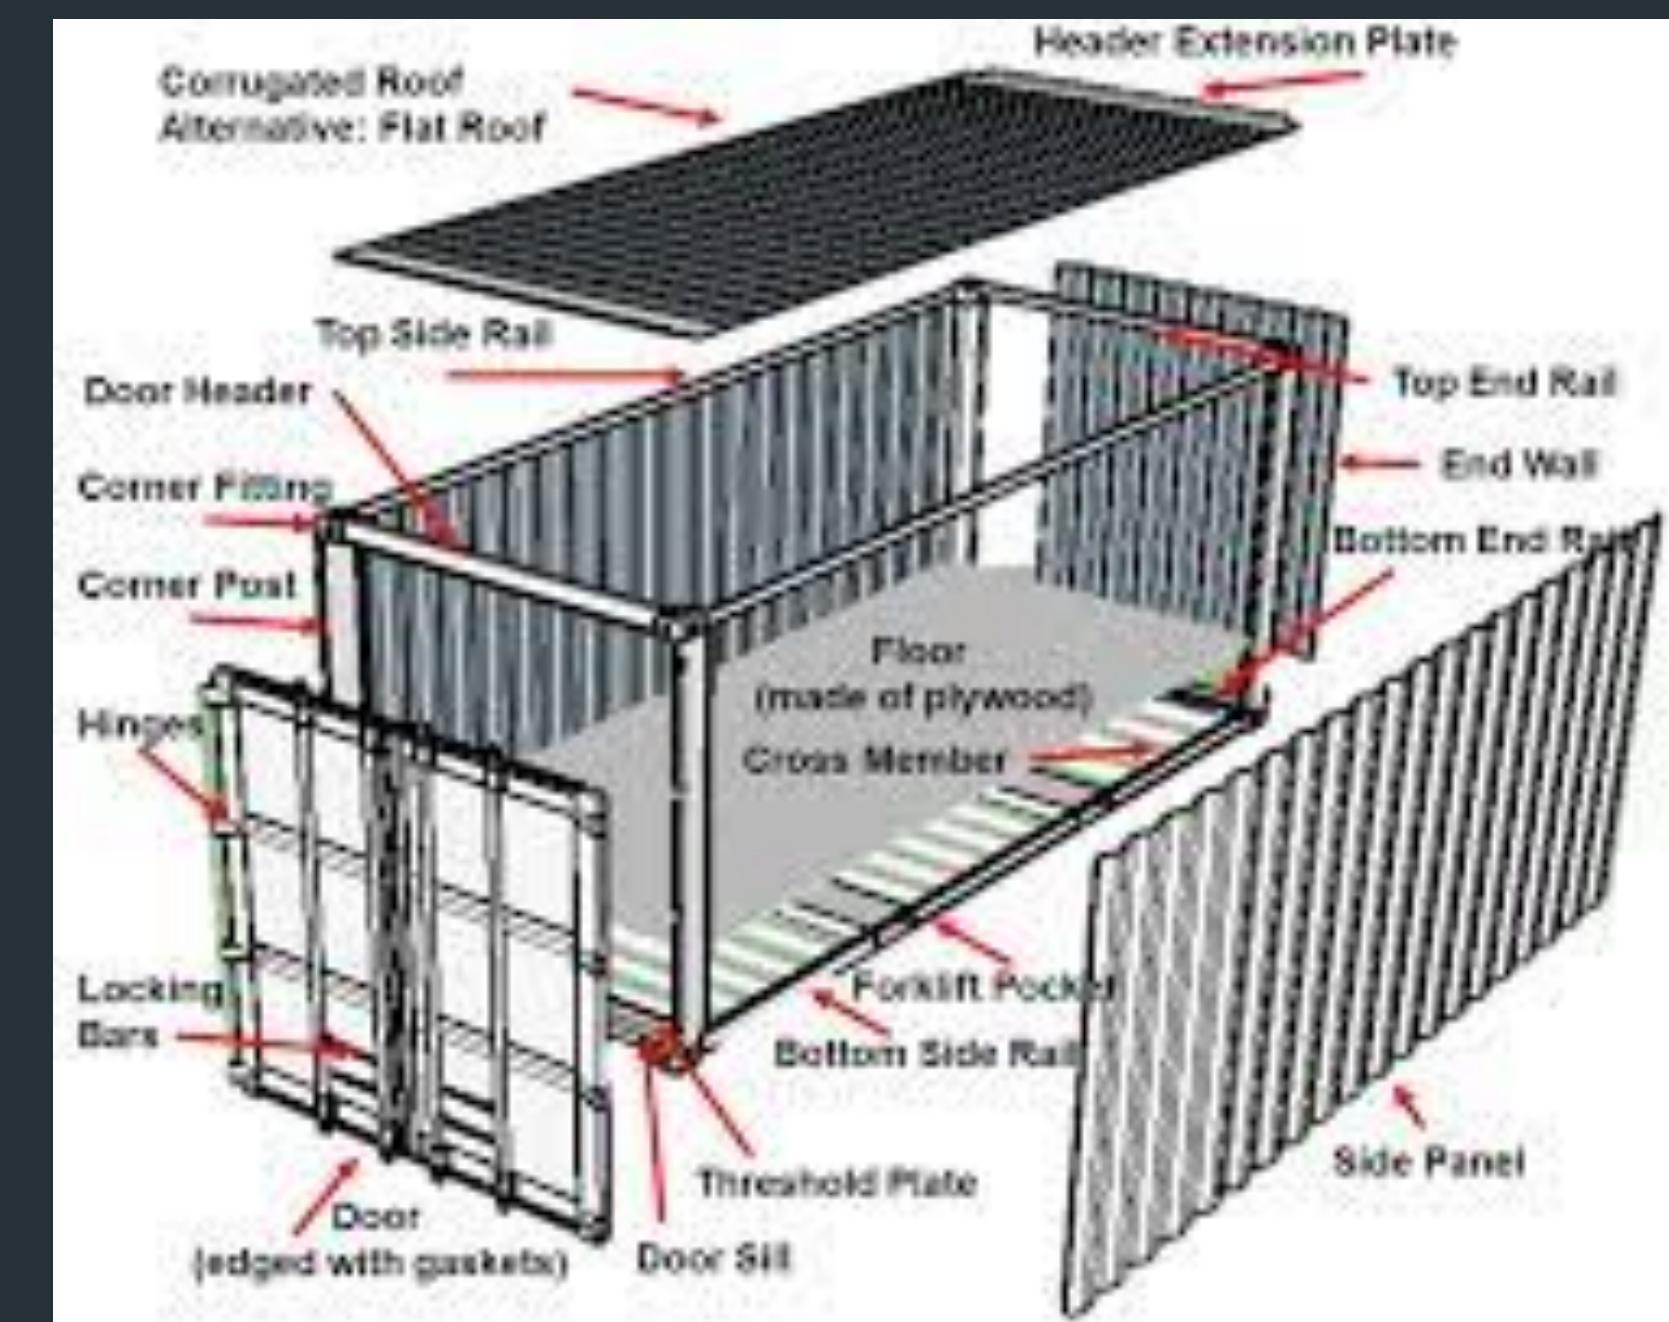
Containers: A Modern Way to Deploy Applications

- Containers are lightweight, portable, and isolated from each other.
- They can be used to deploy applications on a variety of platforms, including servers, clouds, and even on devices.
- Containers are a popular choice for deploying microservices architectures.



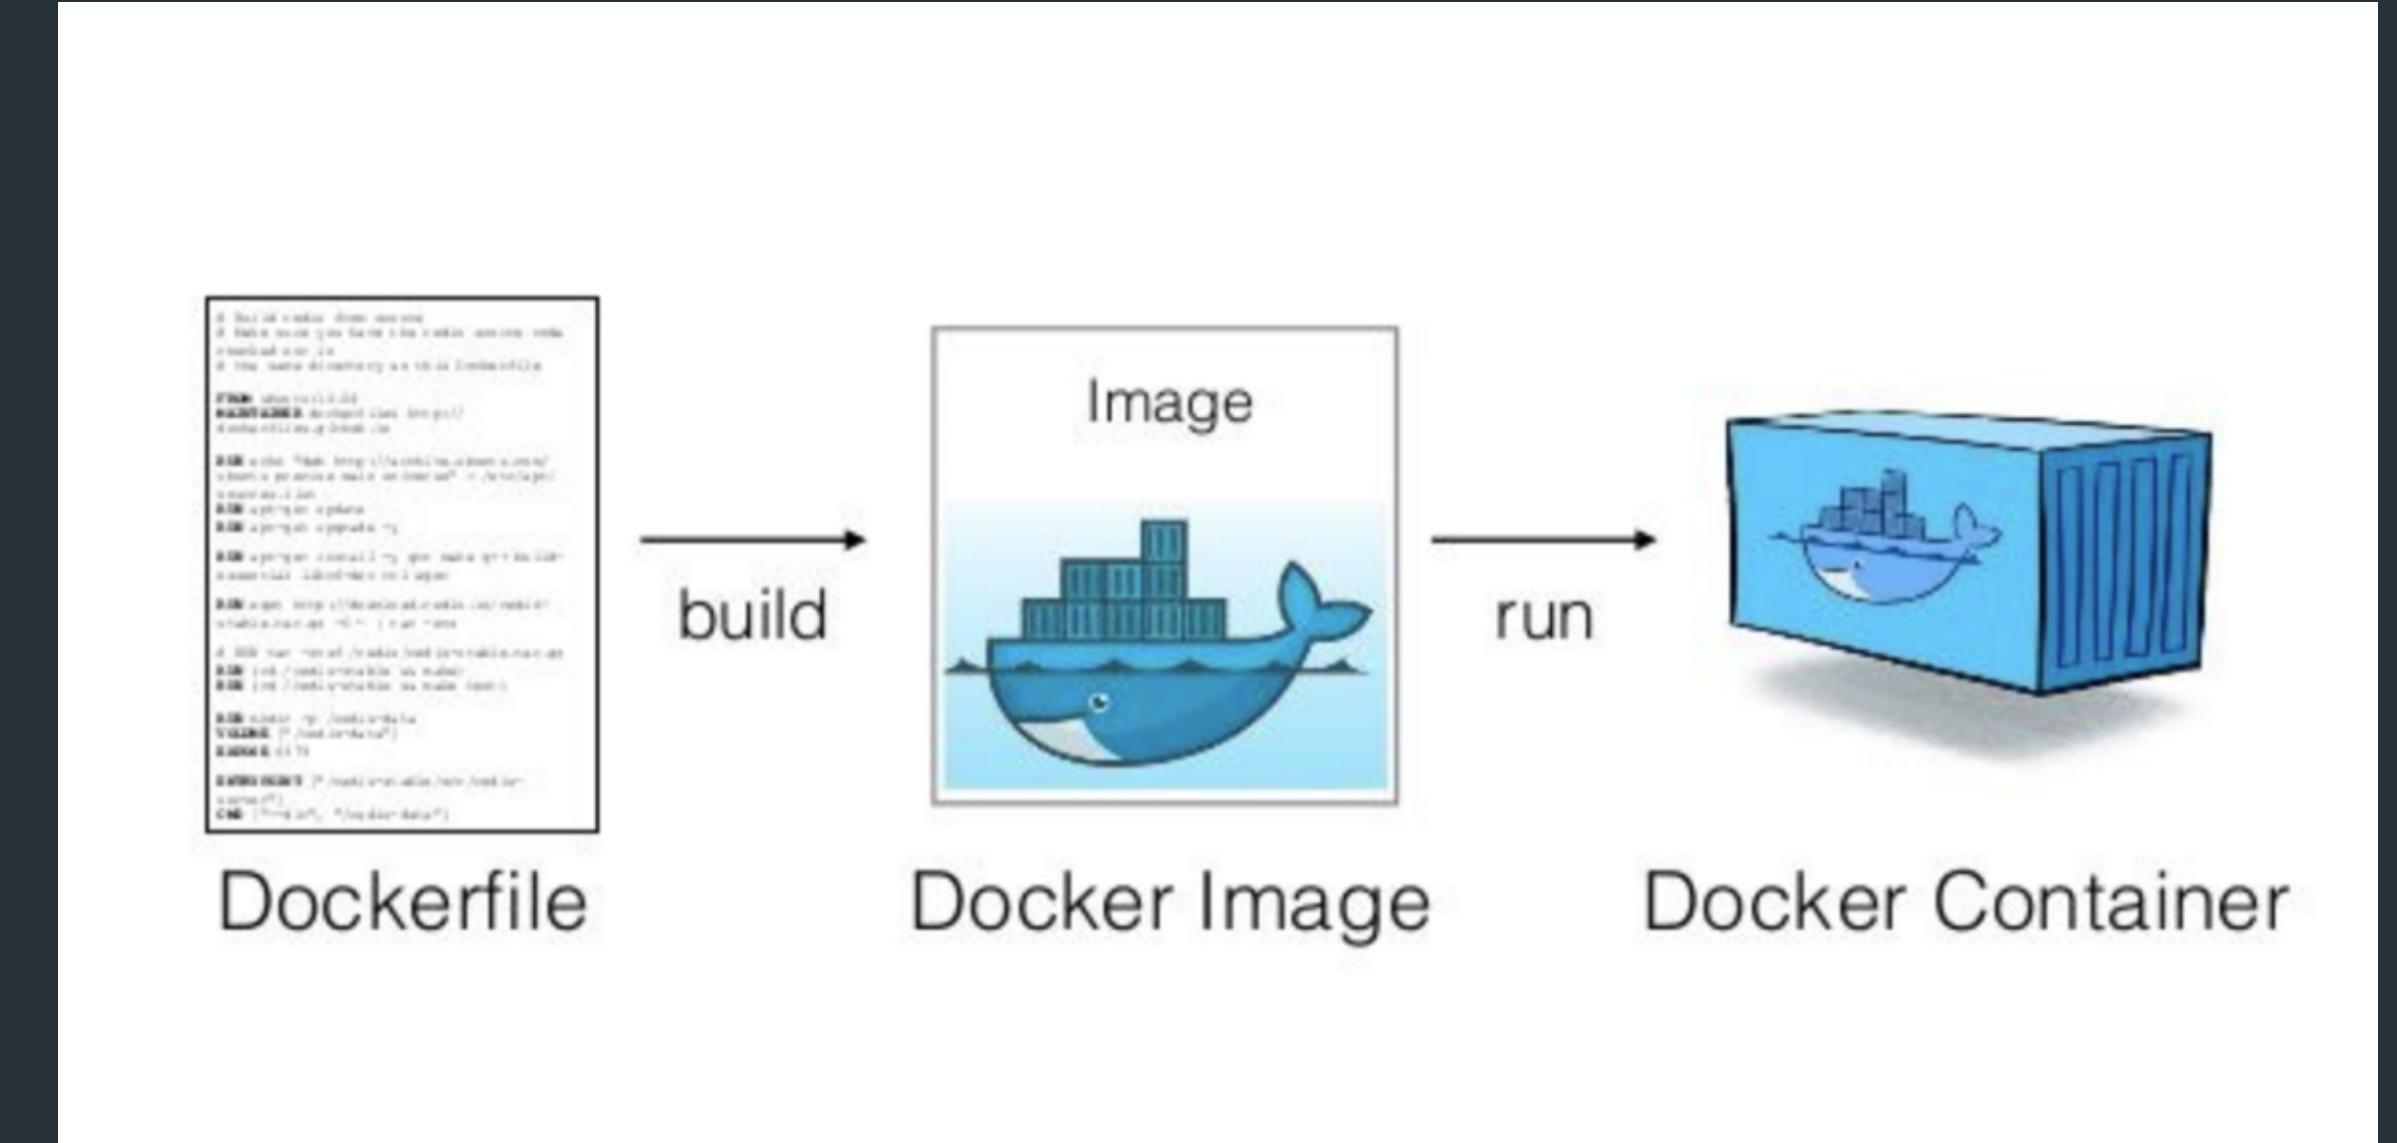
Structure

- **Image:** The image is a read-only template that contains the code, runtime, system tools, system libraries, and settings for a container.
- **Container:** A container is a runnable instance of an image.
- **Container runtime:** The container runtime is responsible for managing the lifecycle of containers.
- **Container orchestration platform:** A container orchestration platform is a tool that helps to manage and scale containerized applications.



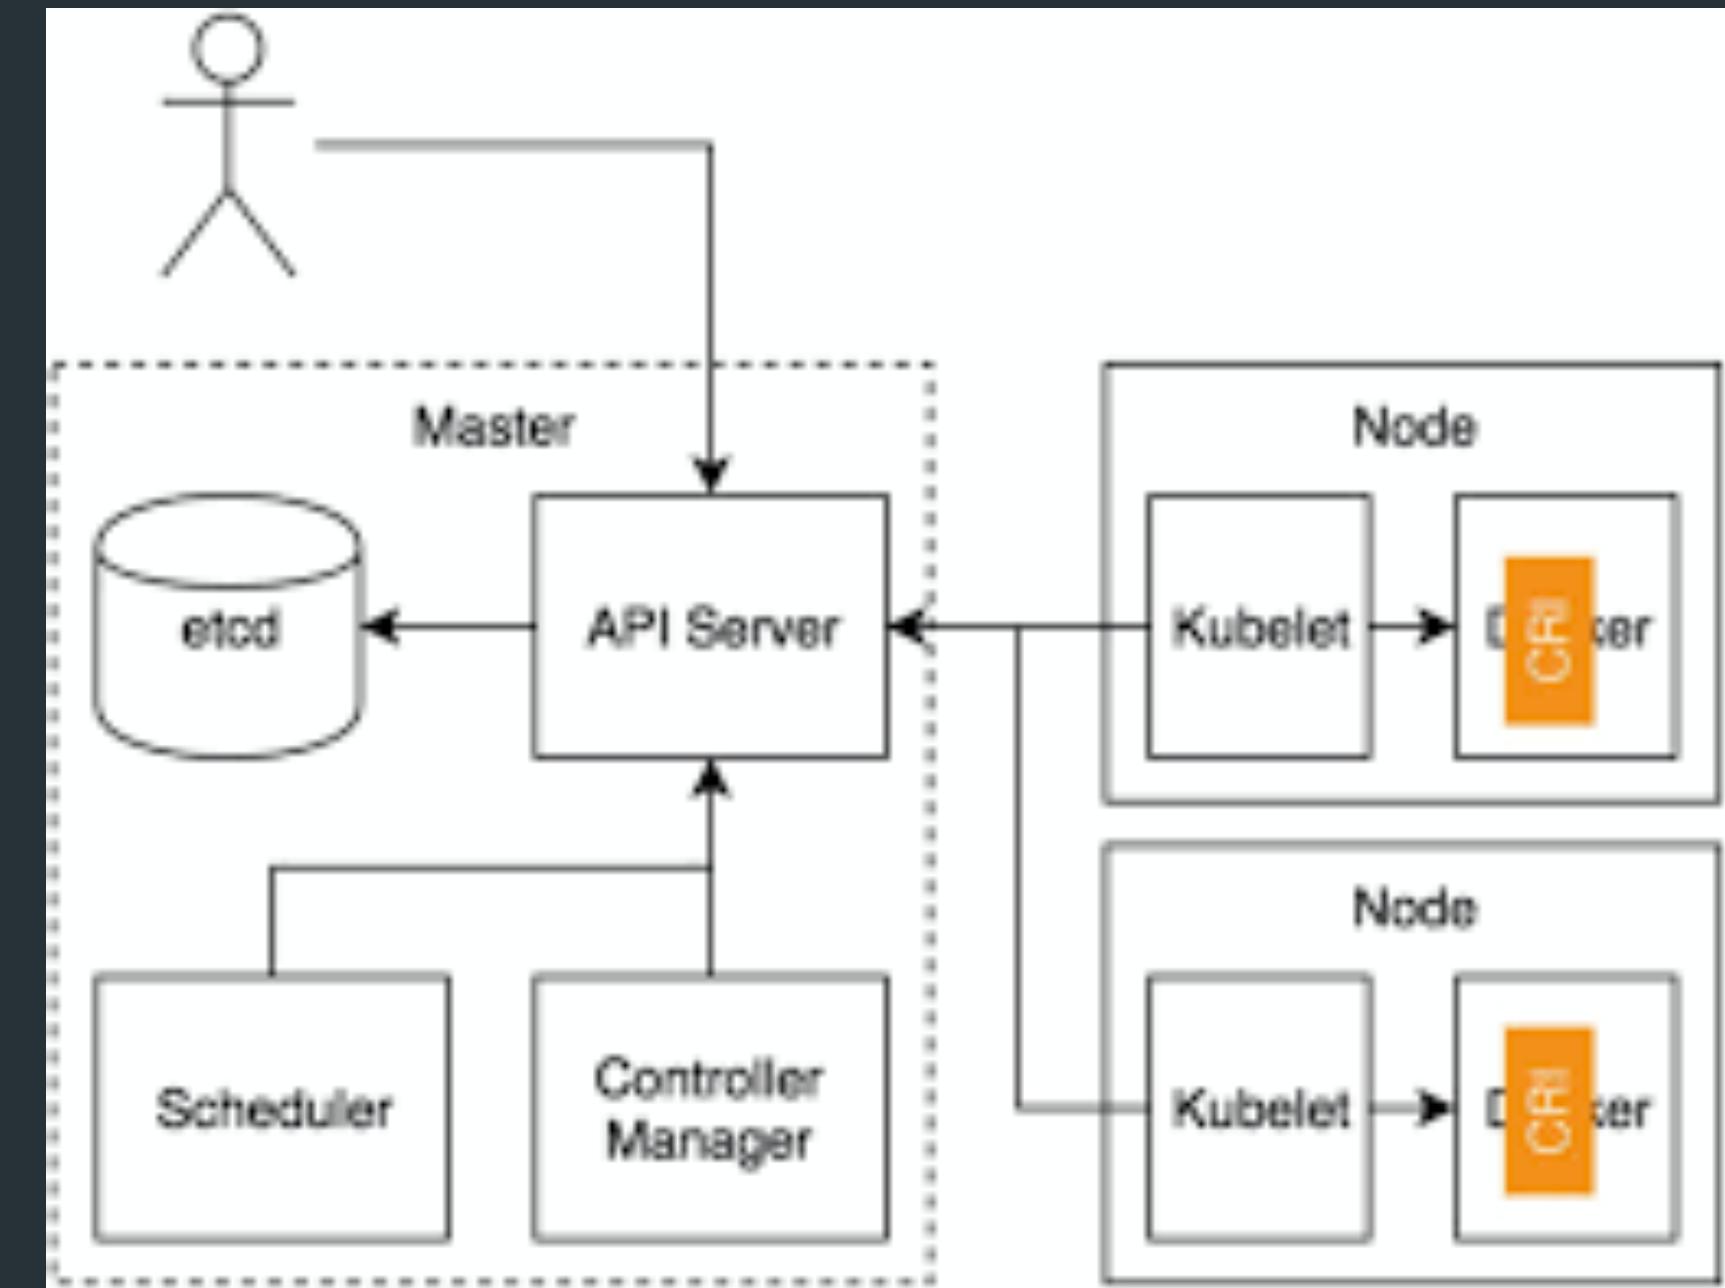
Image

- An image is a read-only template that contains the code, runtime, system tools, system libraries, and settings for a container.
- Images are created using a tool called a Dockerfile.
- Images can be shared and reused.



Runtime

- Responsible for managing the lifecycle of containers.
- Creates, starts, stops, and destroys containers.
- Provides isolation between containers.



Container Orchestration Platform

- A container orchestration platform is a tool that helps to manage and scale containerized applications.
- Container orchestration platforms provide features such as:
 - Automatic deployment
 - Load balancing
 - Autoscaling
 - Monitoring



Objectives

- Containers are a way to package and run applications.
- They are lightweight, portable, and isolated from each other.
- They can be used to deploy applications on a variety of platforms, including servers, clouds, and even on devices.



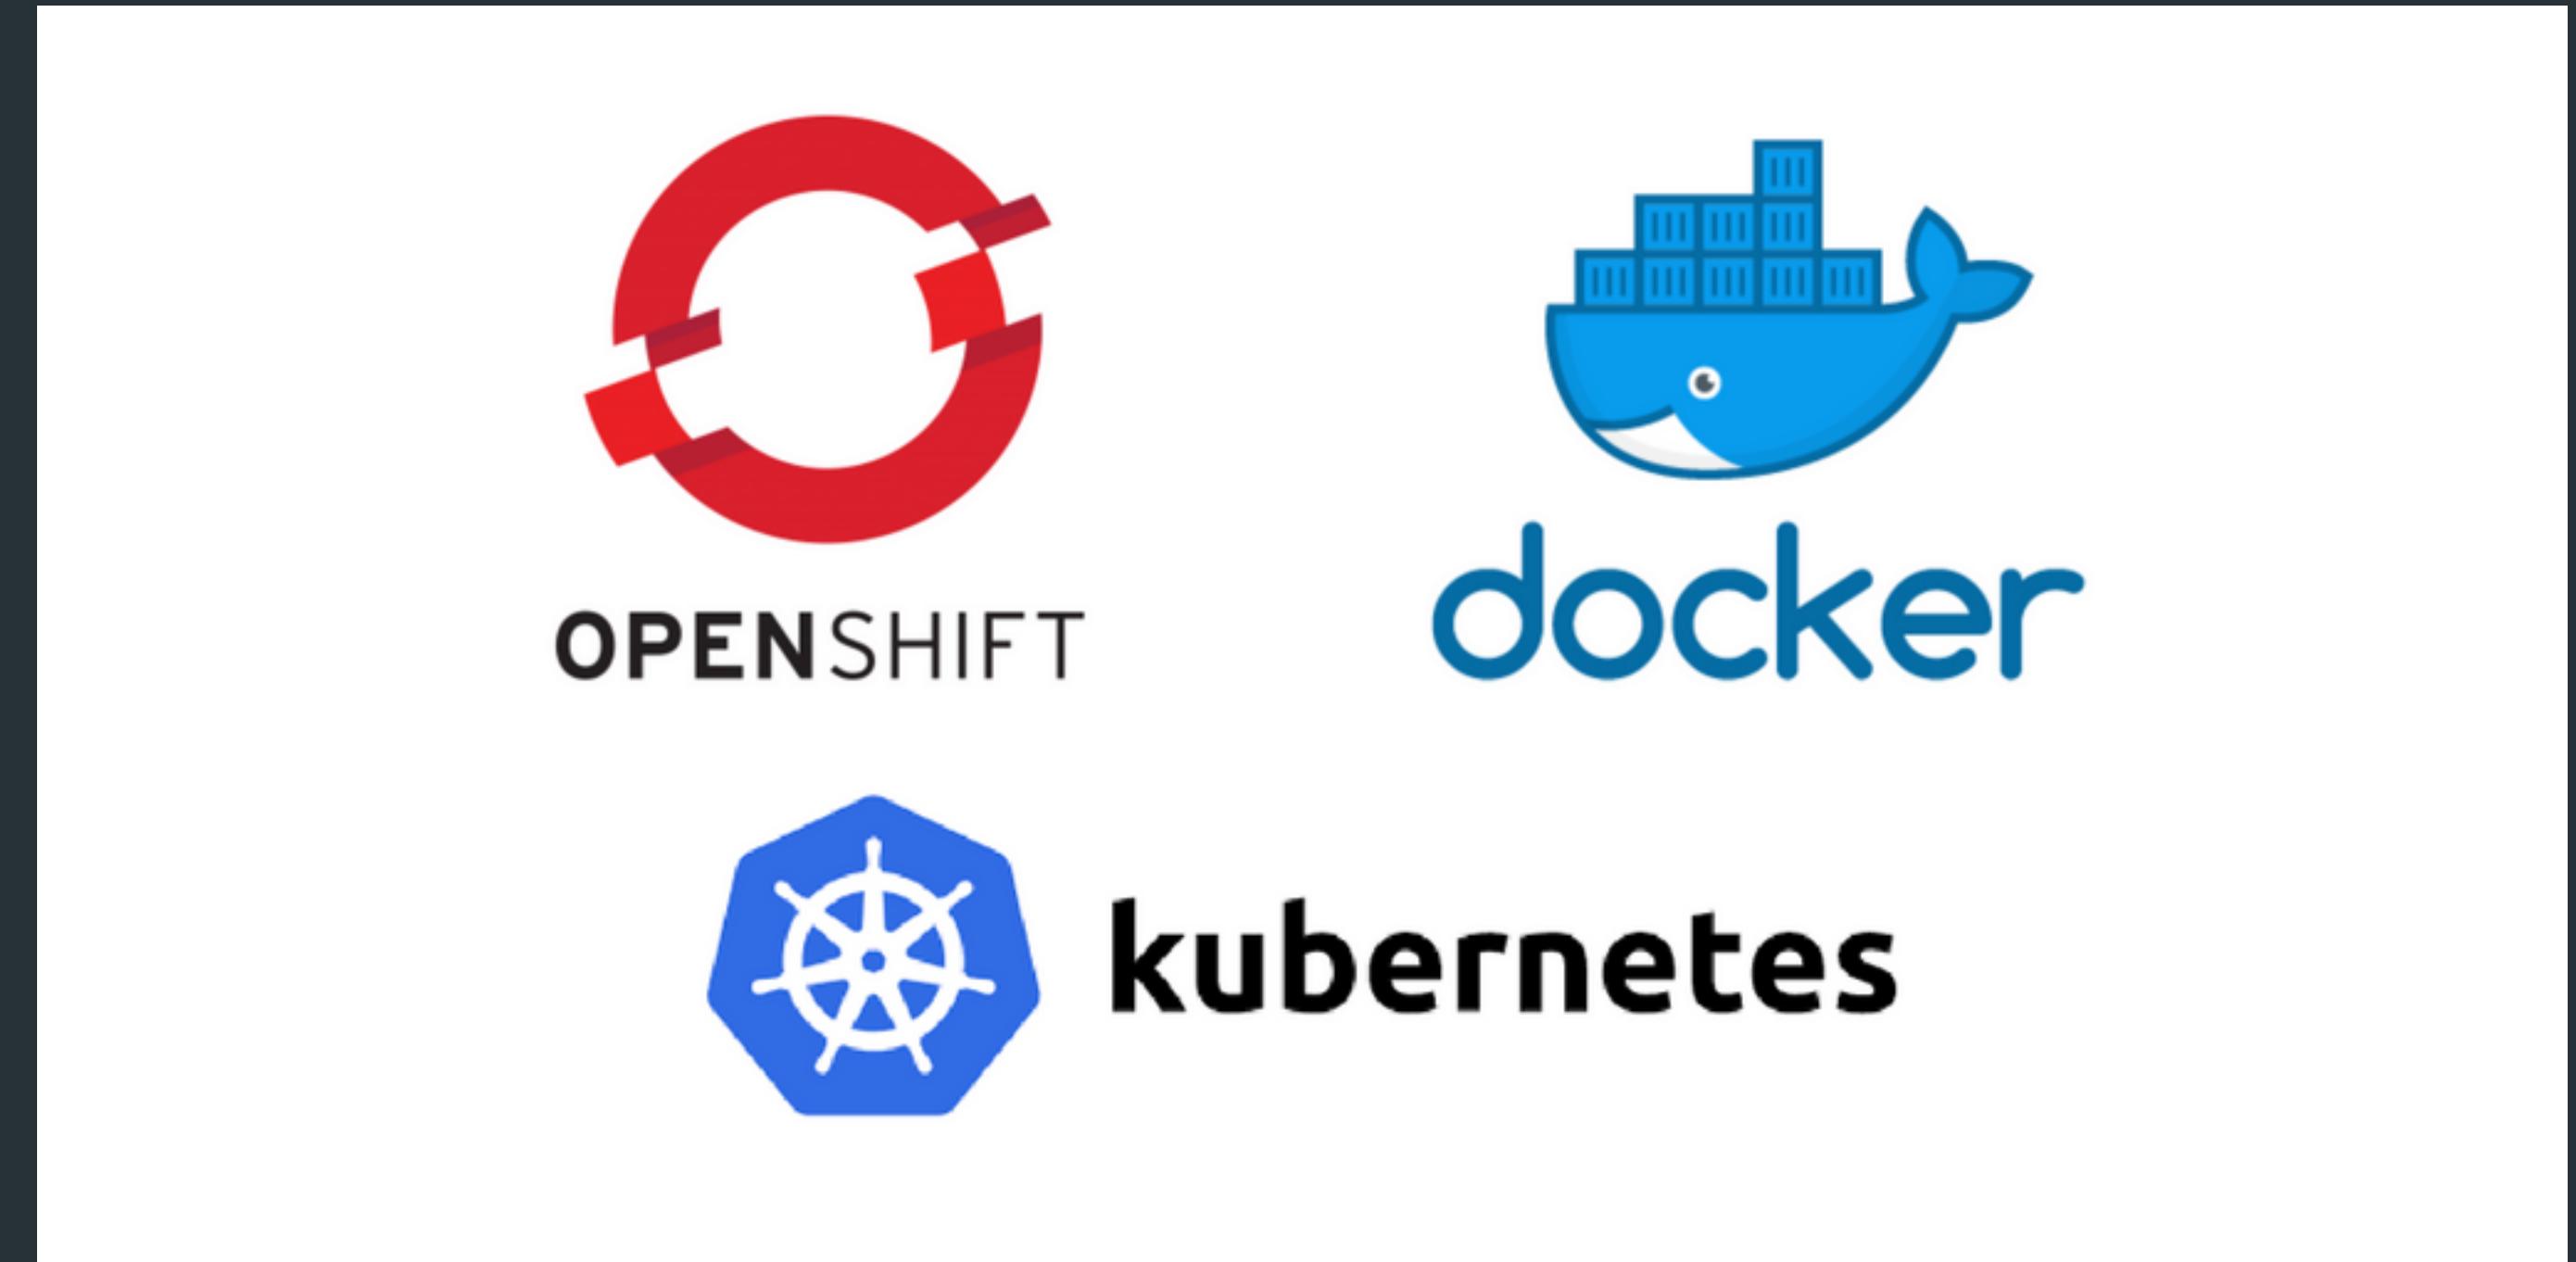
Benefits

- Portability: Containers can be easily moved from one environment to another.
- Isolation: Containers are isolated from each other, which means that they cannot affect each other.
- Efficiency: Containers share the underlying operating system kernel, which makes them more efficient than virtual machines.



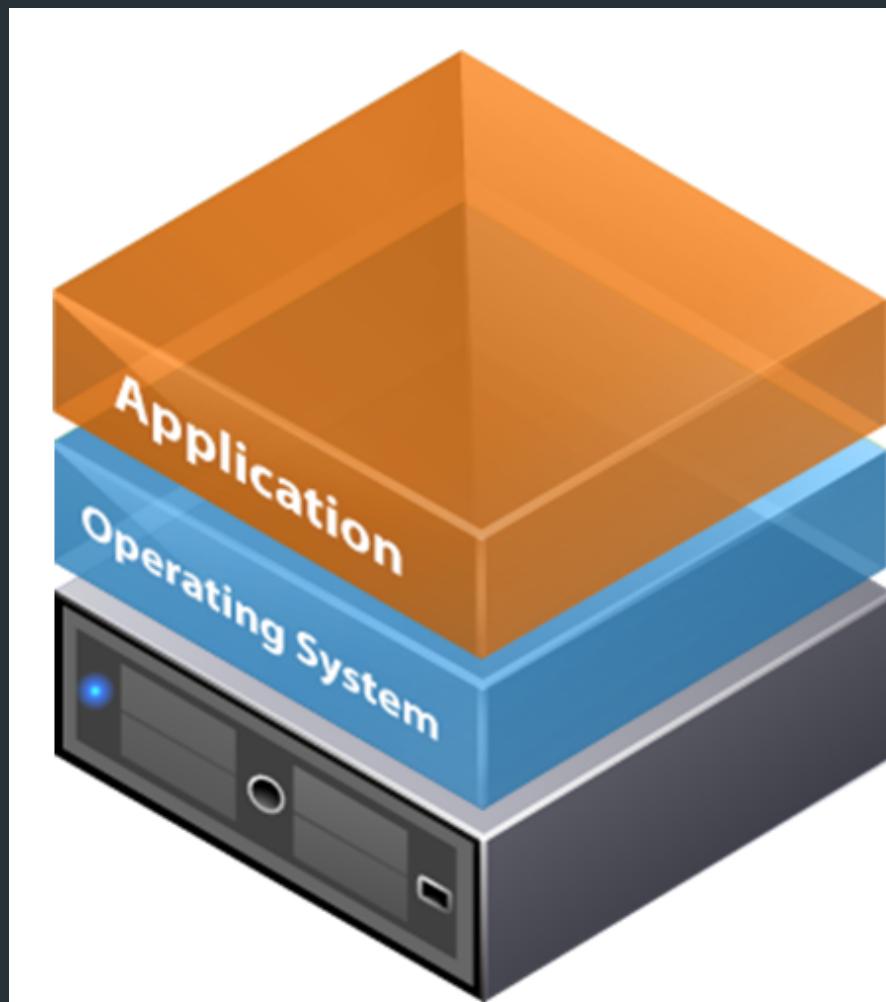
Examples of Containers

- Docker
- Kubernetes
- OpenShift

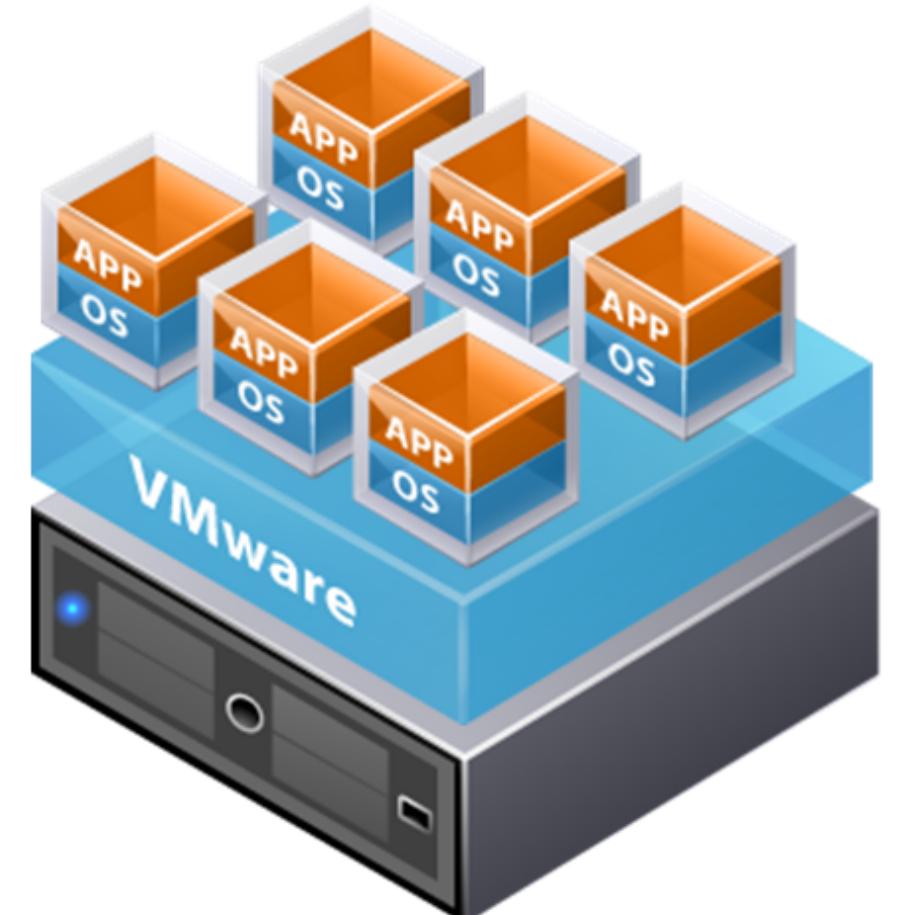


Containers and Virtualization

- Containers are a form of virtualization.
- Containers are more lightweight and portable than virtual machines.



Traditional Architecture



Virtual Architecture

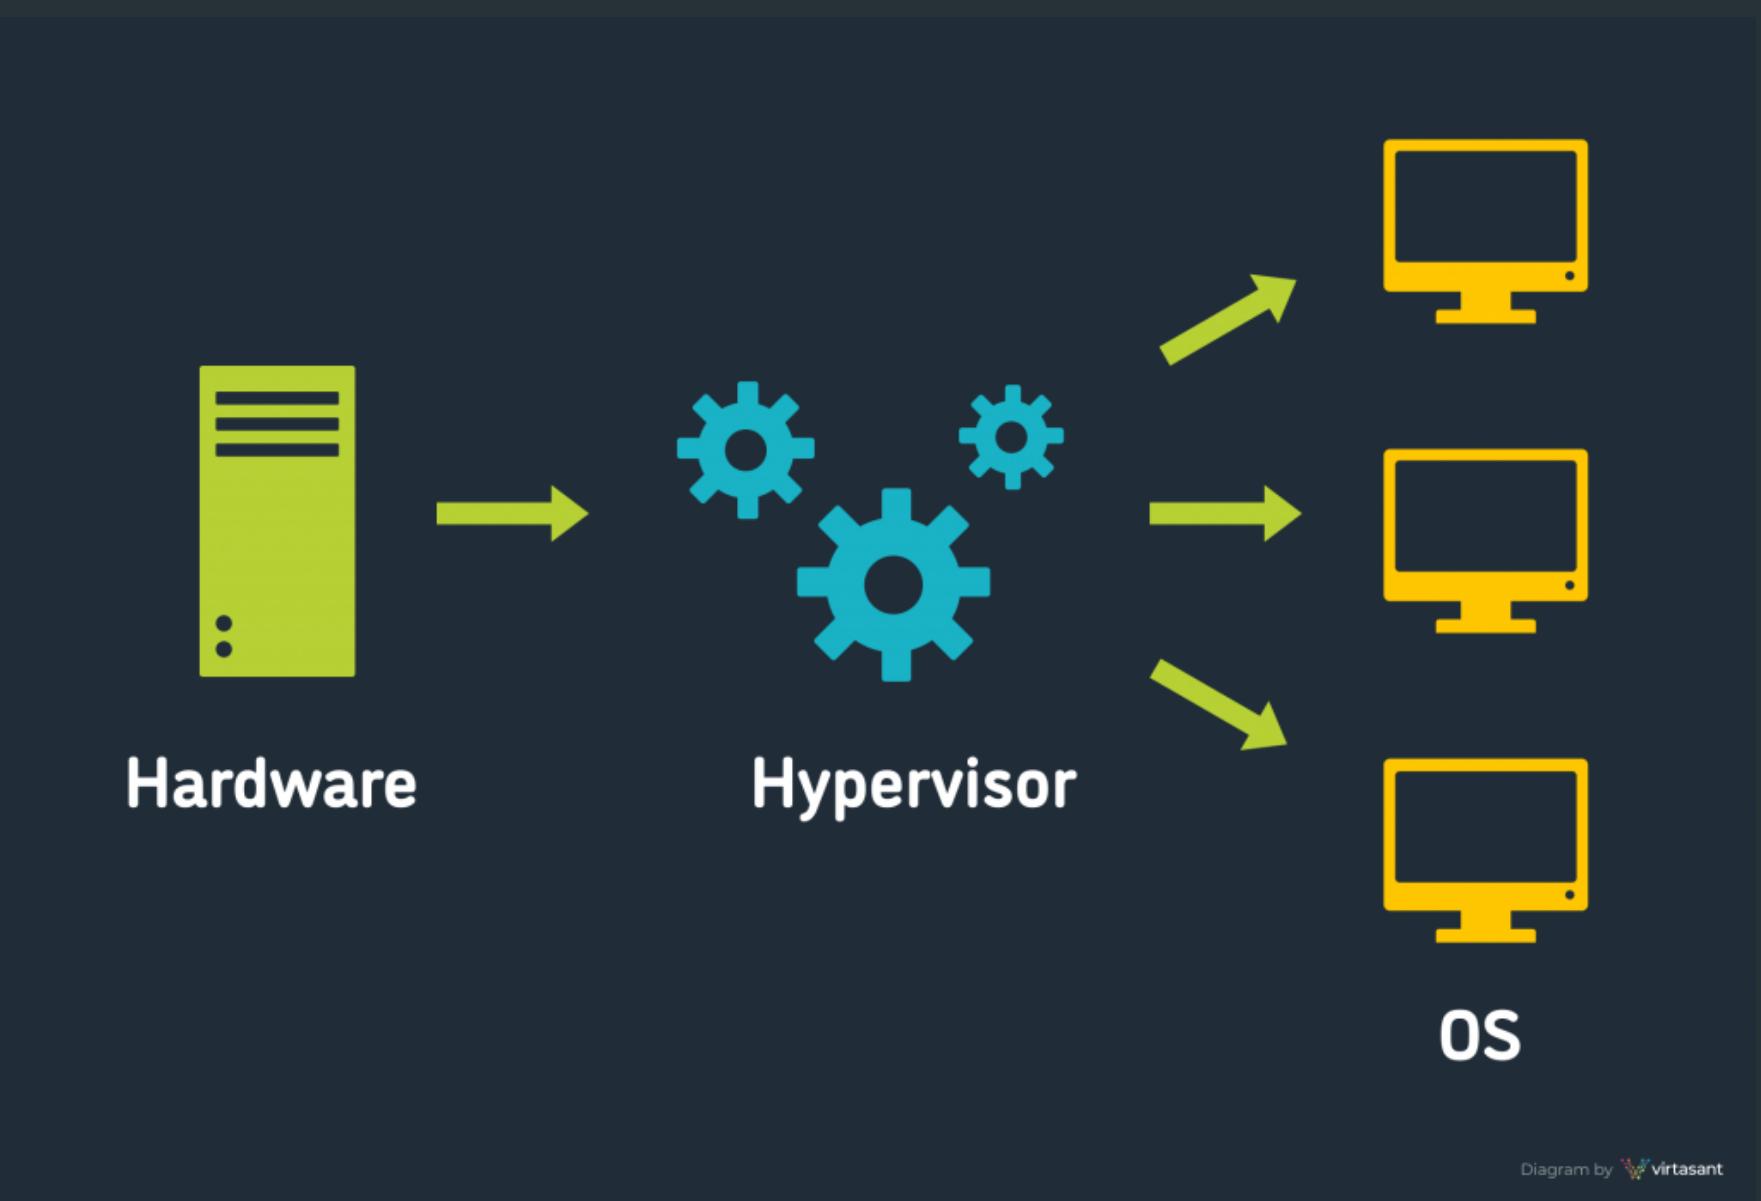
Virtual Machines

- Virtual machines are a way to create isolated operating system environments.
- They are more heavyweight than containers, but they offer more features, such as the ability to run different operating systems on the same physical hardware.
- Virtual machines are a good choice for running legacy applications or applications that require a specific operating system.



Hypervisors

- Hypervisors are the software that allows virtual machines to run on the same physical hardware.
- There are two main types of hypervisors: Type 1 and Type 2.
- Type 1 hypervisors run directly on the physical hardware, while Type 2 hypervisors run on top of an operating system.



Which is right for you?

- Choose containers if you need a lightweight and portable way to run applications.
- Choose virtual machines if you need a more heavyweight solution that offers more features.
- Choose a hypervisor if you need to run multiple virtual machines on the same physical hardware.

YOU'RE
THE ONLY
ONE
WHO
CAN
DECIDE
WHAT'S
BEST
FOR YOU.

DIY: Running a Container on Linux

Step 1: Prepare the System

- Install a Linux distribution that supports containers
- Update the system
- Install the necessary tools
- Create a user group for containers
- Add your user to the container group

```
sudo apt install build-essential  
sudo apt update  
sudo apt install lxc-utils  
sudo groupadd containers  
sudo usermod -aG containers $USER
```

Step 2: Obtaining the Container

- There are two ways to obtain a container image:
- Pull it from a registry
- Build it yourself

```
# Pull an image from a registry  
docker pull nginx
```

```
# Build an image yourself  
docker build -t my-nginx .
```

Step 2b: Build and image yourself

```
FROM nginx:latest

# Copy the nginx configuration file to the container
COPY nginx.conf /etc/nginx/nginx.conf

# Expose port 80 to the host machine
EXPOSE 80

# Start the nginx web server
CMD ["nginx", "-g", "daemon off;"]
```

```
docker build -t my-nginx .
```

Step 3: Setup the group

- Control groups (cgroups) are a Linux kernel feature that allows you to limit and manage resources for groups of processes.
- Cgroups can be used to limit CPU usage, memory usage, disk I/O, and network bandwidth.
- To set up cgroups, you will need to create a cgroup hierarchy.

Create a cgroup hierarchy

```
mkdir -p /cgroup  
mkdir -p /cgroup/cpu  
mkdir -p /cgroup/memory  
mkdir -p /cgroup/disk  
mkdir -p /cgroup/network
```

Step 3: Setup the group

- Control groups (cgroups) are a Linux kernel feature that allows you to limit and manage resources for groups of processes.
- Cgroups can be used to limit CPU usage, memory usage, disk I/O, and network bandwidth.
- To set up cgroups, you will need to create a cgroup hierarchy.

Create a cgroup hierarchy

```
mkdir -p /cgroup  
mkdir -p /cgroup/cpu  
mkdir -p /cgroup/memory  
mkdir -p /cgroup/disk  
mkdir -p /cgroup/network
```

```
cgroups add -g cpu0 nginx
```

Step 3: Setup the group

- Control groups (cgroups) are a Linux kernel feature that allows you to limit and manage resources for groups of processes.
- Cgroups can be used to limit CPU usage, memory usage, disk I/O, and network bandwidth.
- To set up cgroups, you will need to create a cgroup hierarchy.

Create a cgroup hierarchy

```
mkdir -p /cgroup  
mkdir -p /cgroup/cpu  
mkdir -p /cgroup/memory  
mkdir -p /cgroup/disk  
mkdir -p /cgroup/network
```

```
cgroups add -g cpu0 nginx
```

```
cgroups set -r cpu.shares=50 nginx
```

```
cgroups stat -c cpu nginx
```

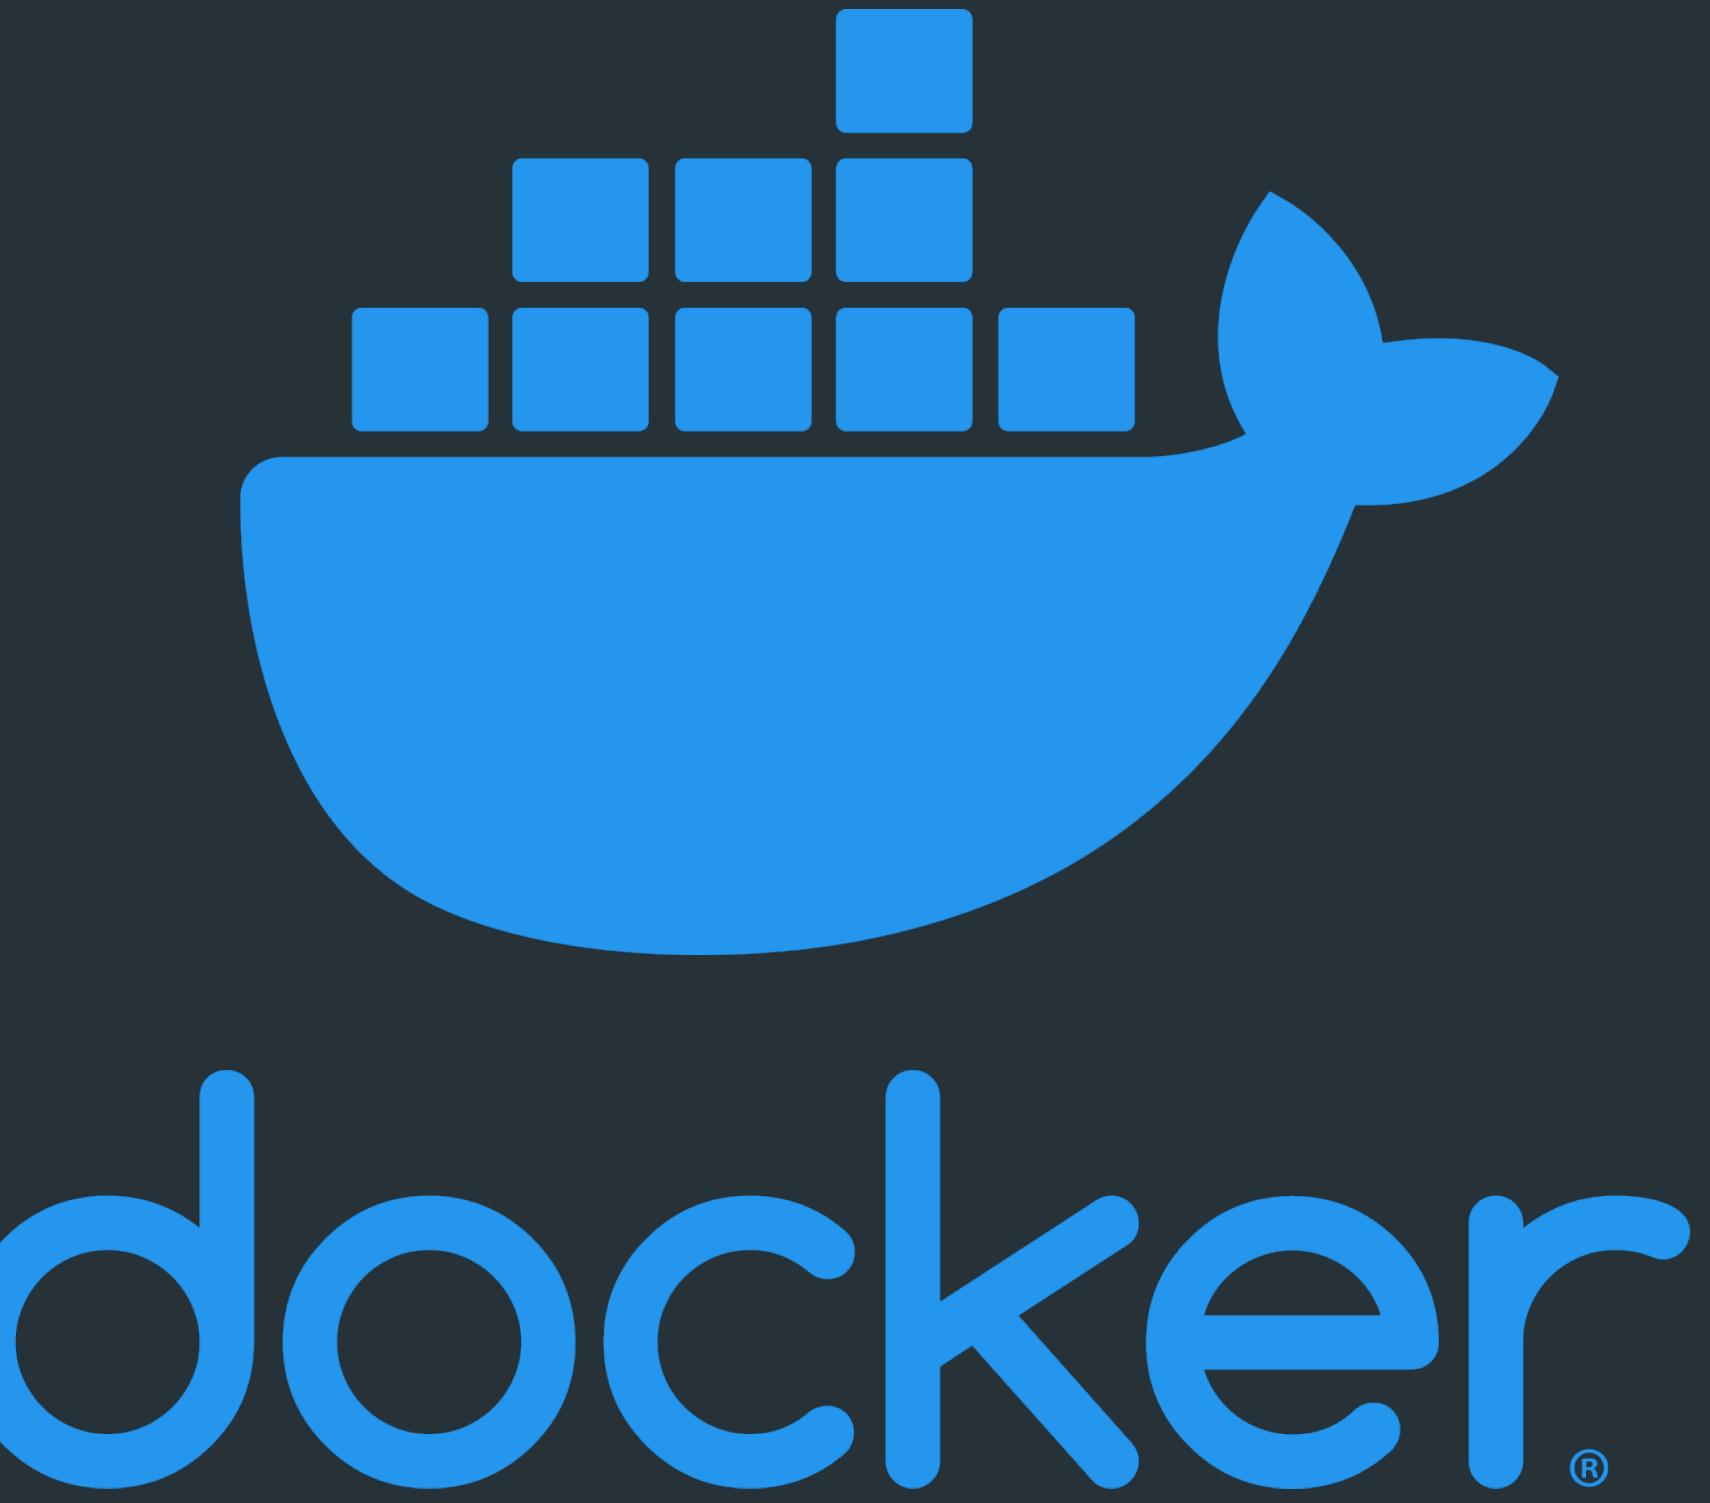
Step 3: Run the Container

- To run a container, you will need to use the docker run command.
- The docker run command takes a number of arguments, including the name of the image to run, the command to run inside the container, and the environment variables to set.
- For example, to run the nginx web server on port 80, you would use the following command:

```
docker run -p 80:80 nginx
```

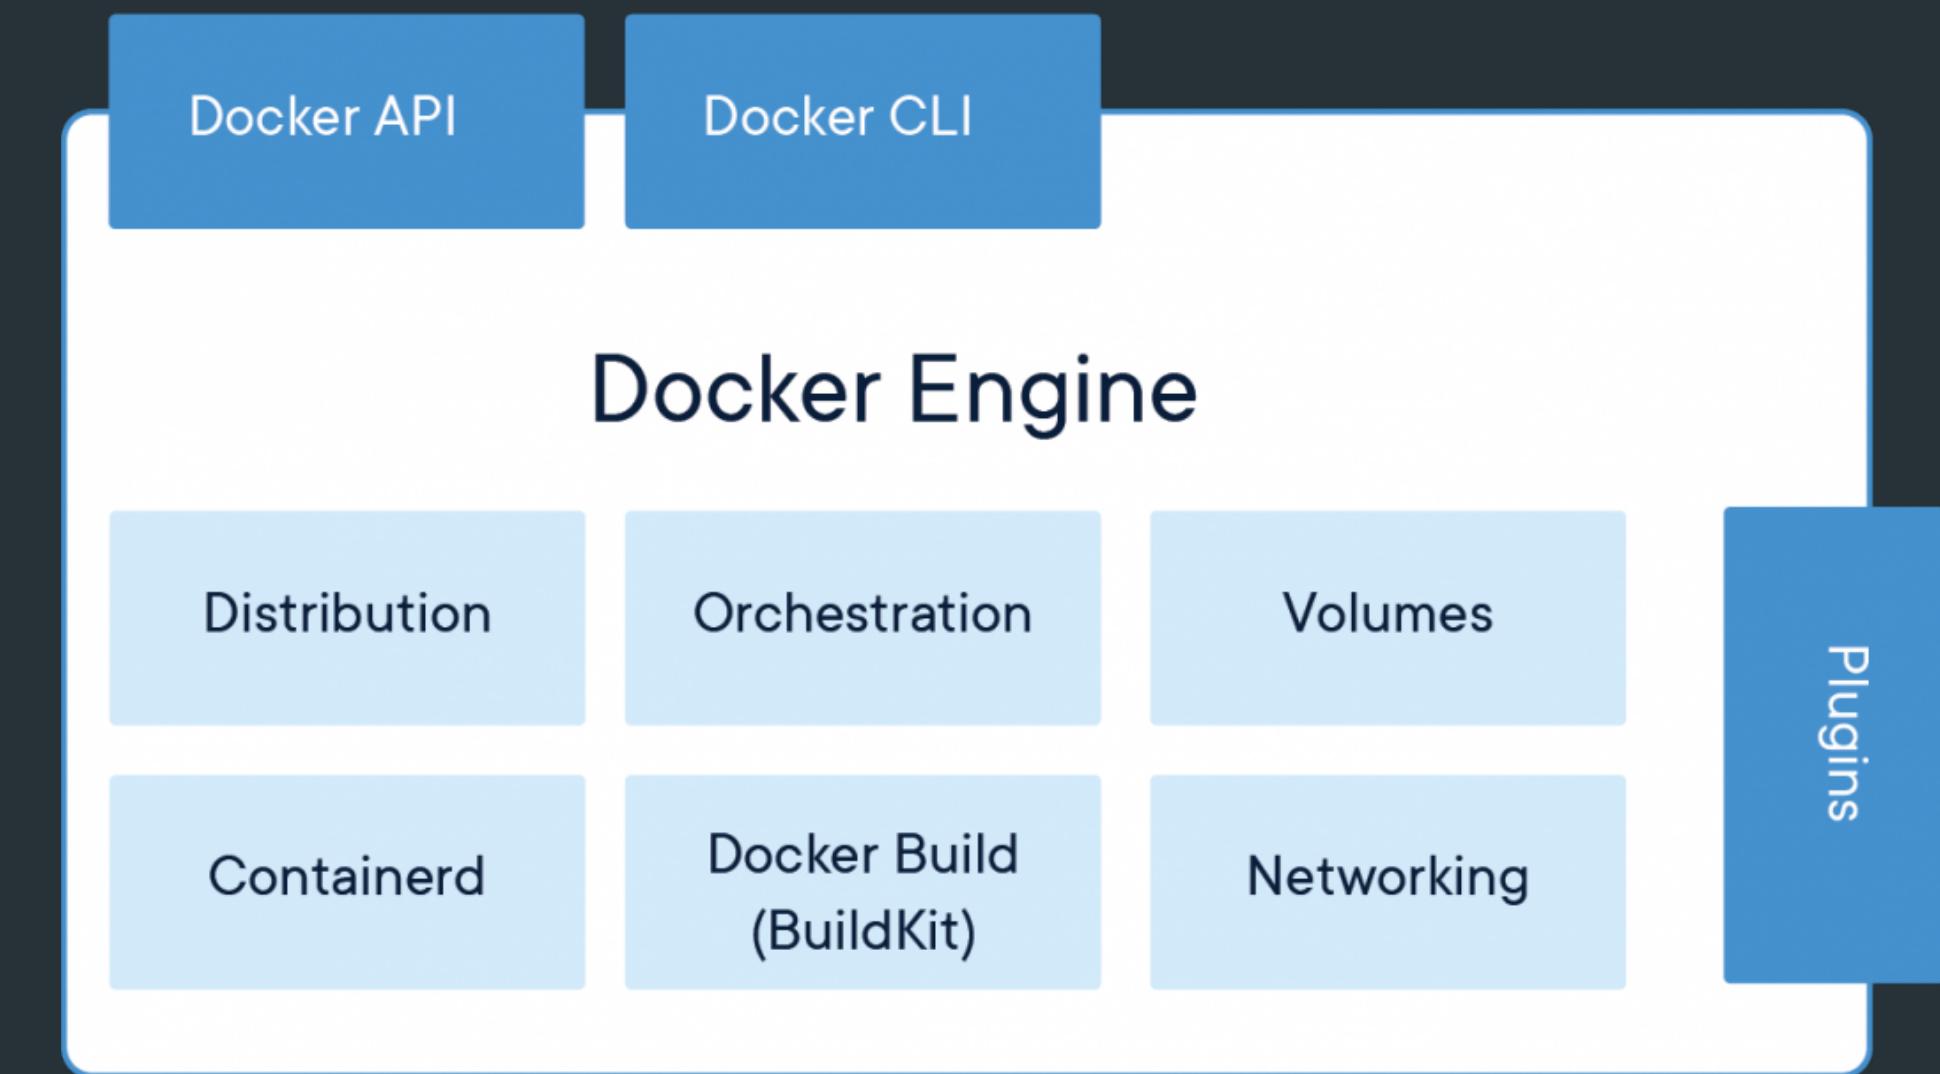
Docker

- Engine
- Daemon
- Registry
- Client



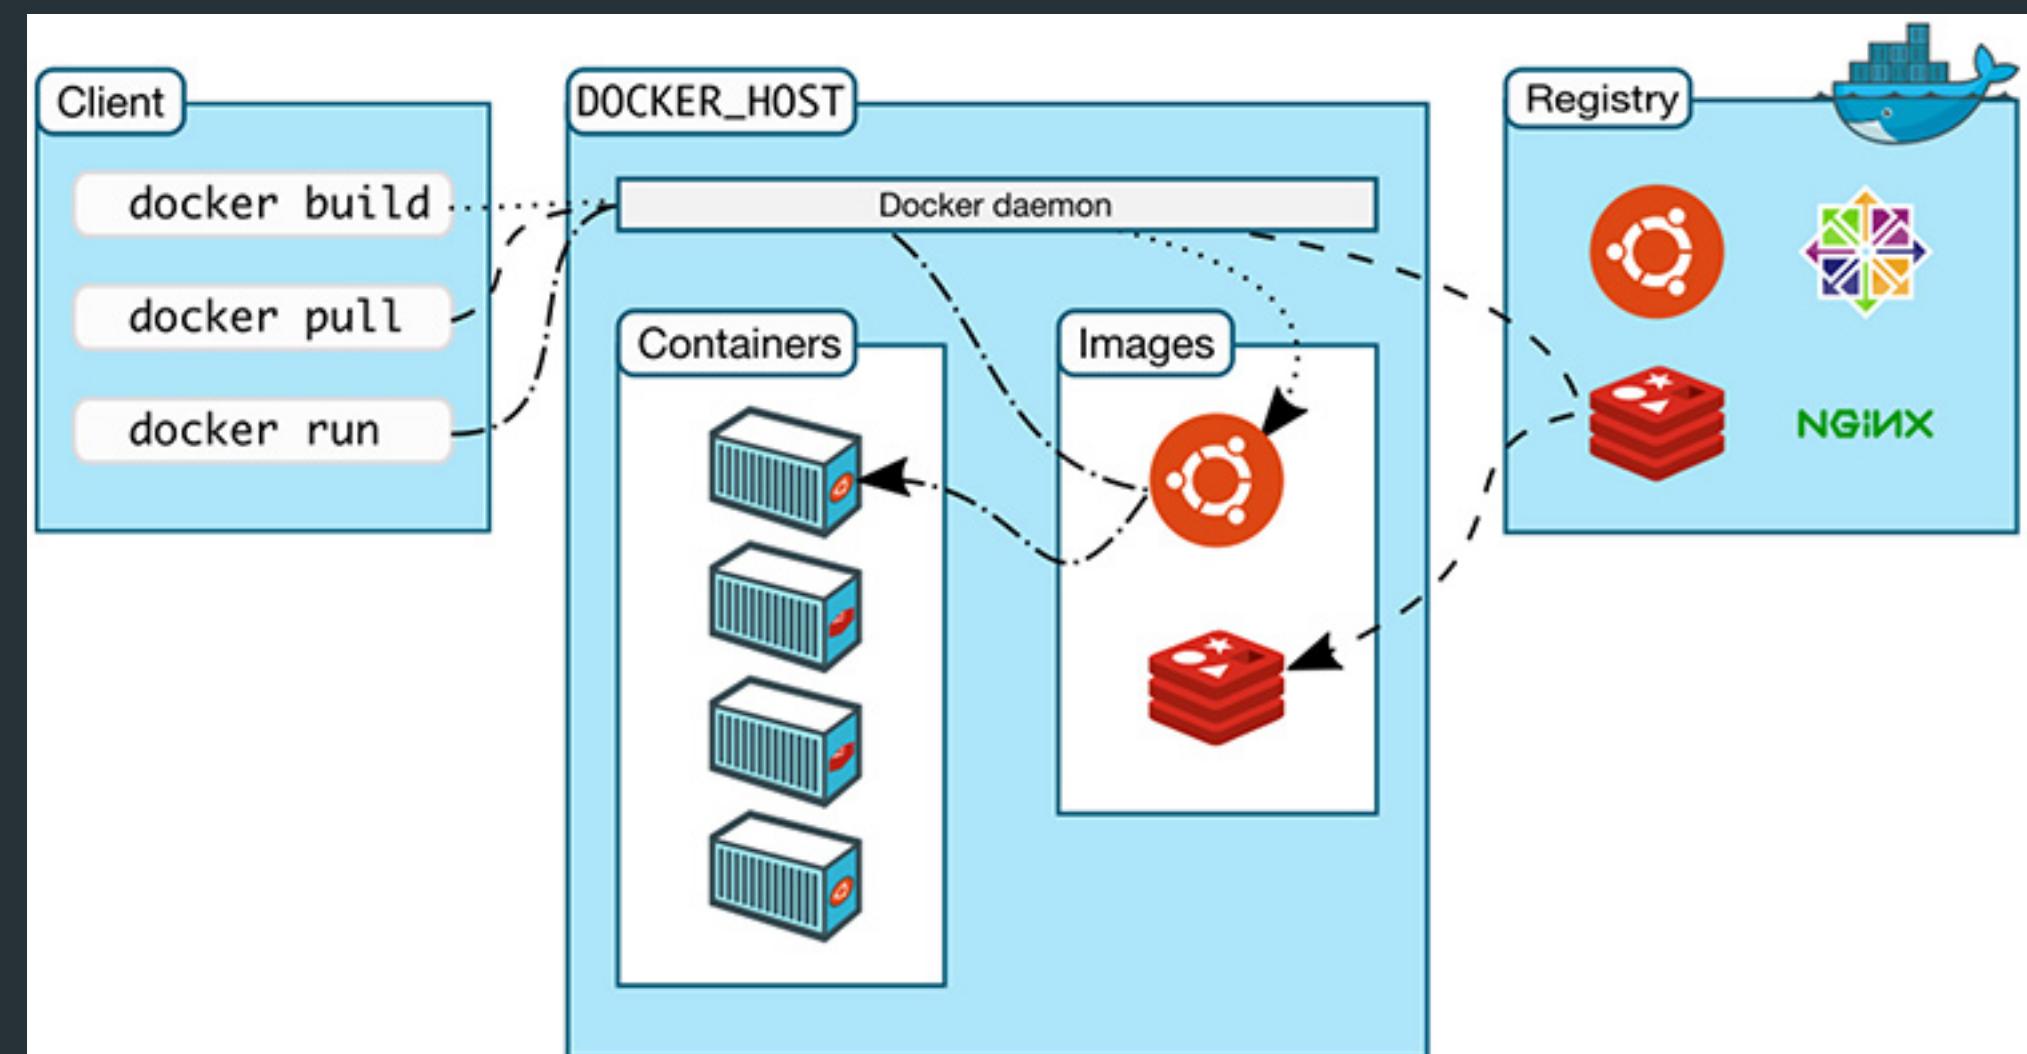
Docker Engine

- Is the core component of Docker that manages containers.
- Responsible for creating, starting, stopping, and removing containers.
- Manages the resources that are allocated to containers.



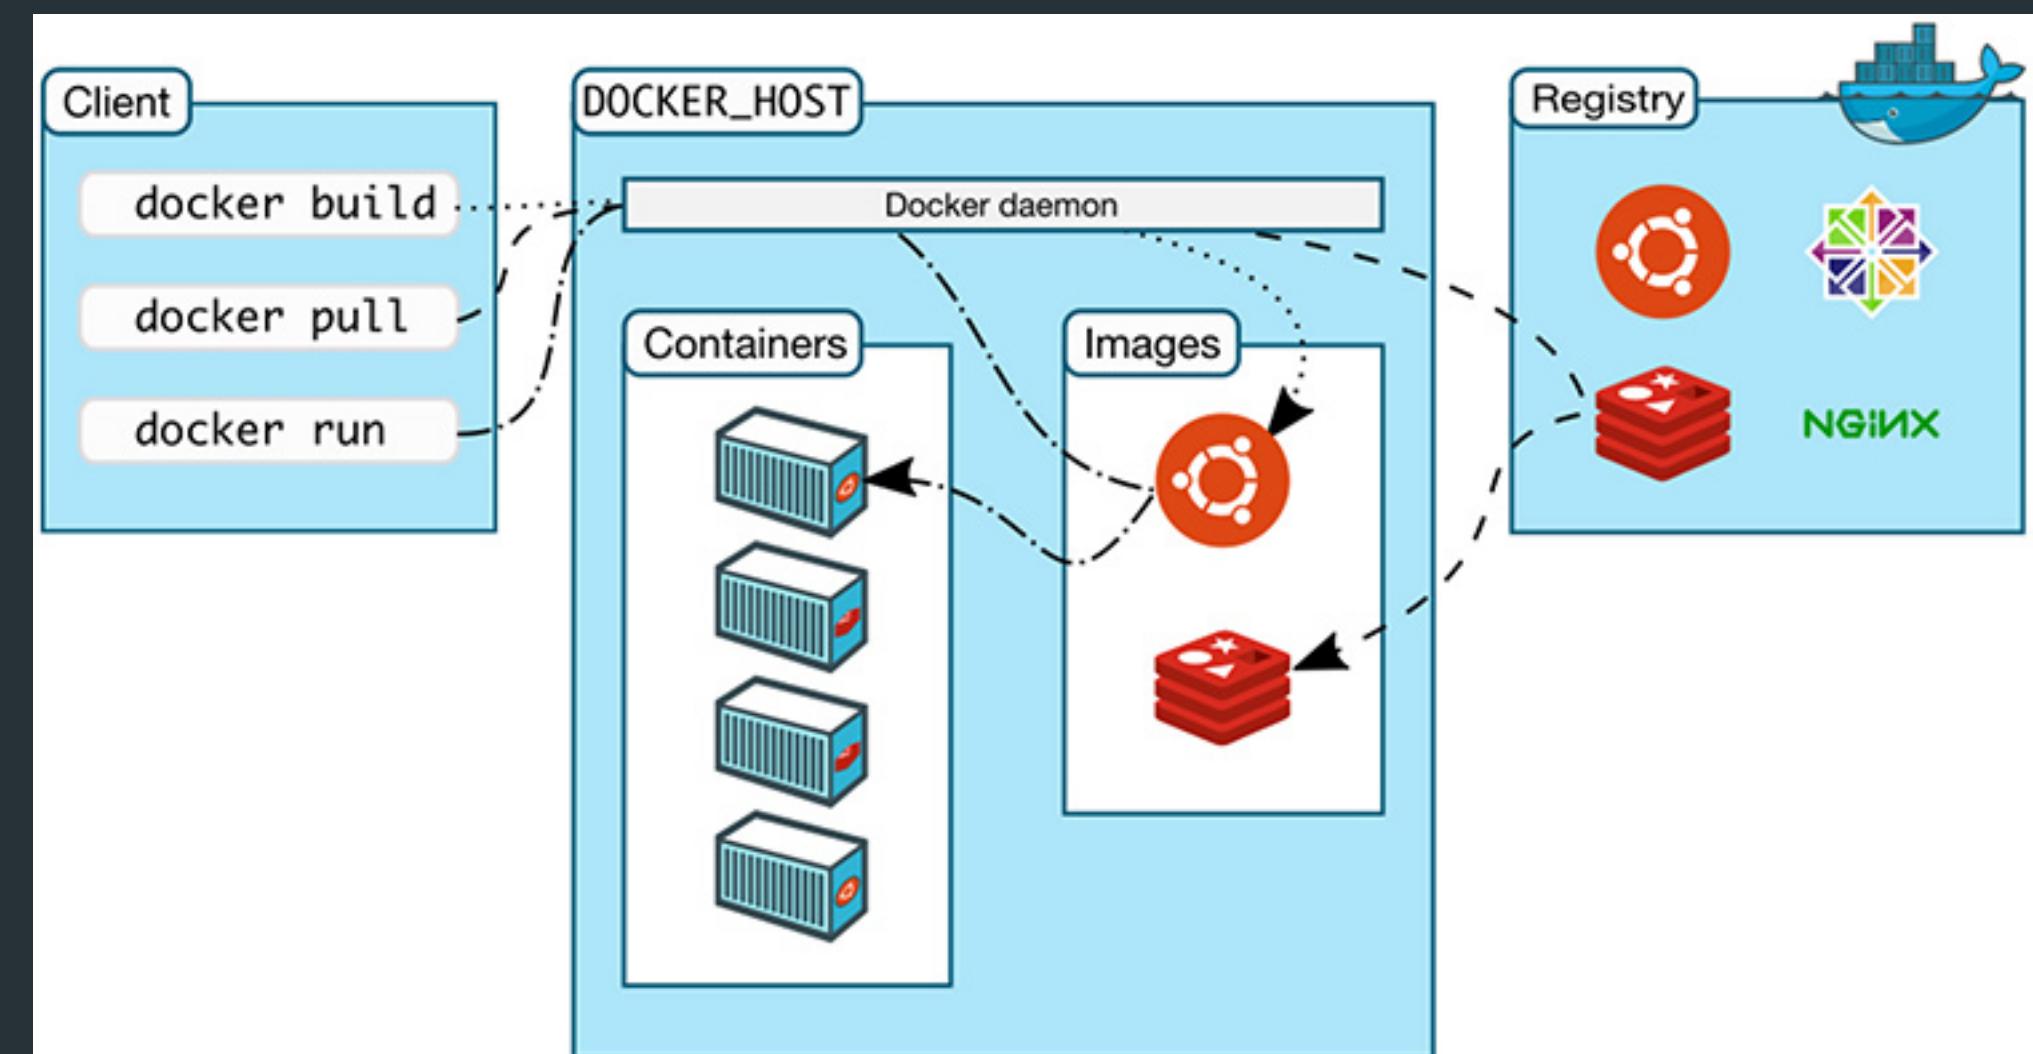
Docker Daemon

- A background process that runs on the host machine and manages the Docker Engine.
- Responsible for communicating with the Docker Registry and downloading Docker images.
- Manages the containers that are running on the host machine.



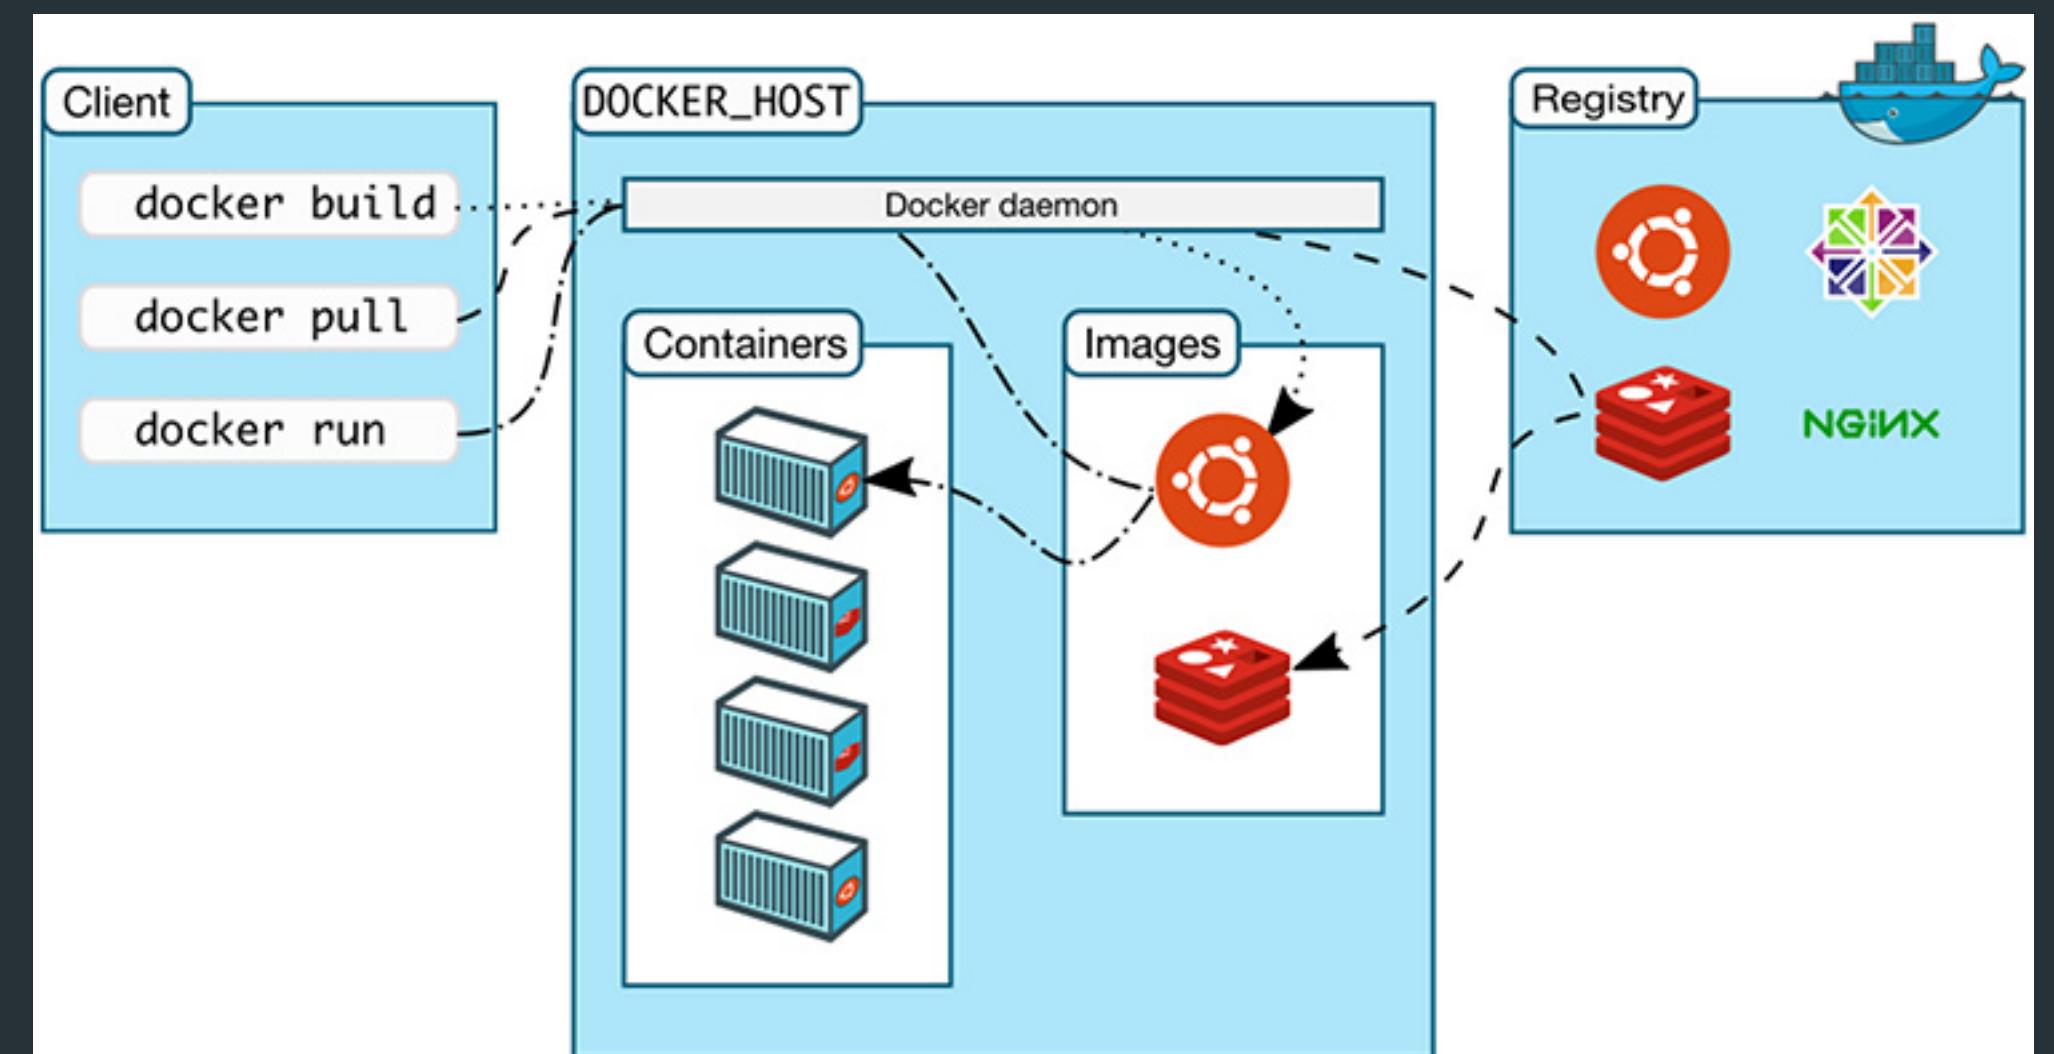
Docker Registry

- A central repository where Docker images can be stored and shared.
- Docker images are essentially read-only templates that can be used to create containers.
- Docker images can be created from scratch or they can be built from other Docker images.



Docker Client

- A command-line tool that allows you to interact with the Docker Engine.
- Used to create, start, stop, and remove containers.
- Used to manage the resources that are allocated to containers.



Writing Docker Files

- A Docker file is a text file that contains the instructions for building a Docker image.
- Docker files are written in a simple, human-readable format.
- Docker files can be used to build Docker images for any application.

```
FROM nginx:latest
```

```
COPY ./usr/share/nginx/html
```

```
EXPOSE 80
```

```
CMD ["nginx", "-g", "daemon off;"]
```

Writing Docker Files

- A typical Docker file will include the following steps:

- FROM: This specifies the base image that the Docker image will be built on.

FROM nginx:latest

- RUN: This specifies commands that will be run when the Docker image is built.

COPY ./usr/share/nginx/html

- COPY: This specifies files or directories that will be copied into the Docker image.

EXPOSE 80

- EXPOSE: This specifies ports that will be exposed by the Docker image.

CMD ["nginx", "-g", "daemon off;"]

- CMD: This specifies the command that will be run when the Docker image is started.

More Docker Instructions

- ARG: Defines an environment variable that can be used in the Docker file.
- AND: Specifies that the following commands should be run only if the previous command succeeds.
- ENTRYPOINT: The command that will be run when the Docker image is started.
- WORKDIR: The working directory for the Docker image.
- STOPSIGNAL: The signal that will be sent to the Docker container when it is stopped.
- USER: The user that will run the commands in the Docker image.
- LABEL: A label for the Docker image.
- ENV: An environment variable for the Docker image.

`ARG VERSION=1.0`

`RUN apt-get update && apt-get install -y nginx`

`ENTRYPOINT nginx`

`WORKDIR /usr/share/nginx/html`

`STOPSIGNAL SIGTERM`

`USER nginx`

`LABEL app=nginx`

`ENV PORT=80`

Docker Commands

- Some of the most common Docker commands include:
 - docker build: Builds a Docker image from a Dockerfile.
 - docker run: Starts a Docker container from an image.
 - docker stop: Stops a Docker container.
 - docker remove: Removes a Docker container.
 - docker inspect: Inspects a Docker container.
 - docker logs: Displays the logs for a Docker container.

Build a Docker image from a Dockerfile
docker build -t my-app .

Start a Docker container from an image
docker run -p 80:80 my-app

Stop a Docker container
docker stop my-app

Remove a Docker container
docker rm my-app

Inspect a Docker container
docker inspect my-app

Display the logs for a Docker container
docker logs my-app

More Docker Commands

List all containers

```
docker ls
```

Show the history of a container

```
docker history my-container
```

Search for Docker images

```
docker search nginx
```

Pull an image from a registry

```
docker pull nginx
```

Tag an image

```
docker tag nginx:latest my-nginx
```

Push an image to a registry

```
docker push my-nginx
```

Rename a container

```
docker rename my-container my-new-container
```

Attach to a running container

```
docker attach my-container
```

```
# Tag an image  
docker tag nginx:latest my-nginx
```

```
# Push an image to a registry  
docker push my-nginx
```

```
# Rename a container  
docker rename my-container my-new-container
```

```
# Attach to a running container  
docker attach my-container
```

```
# Commit a container's changes to an image  
docker commit my-container my-new-image
```

```
# Get statistics for a container  
docker stats my-container
```

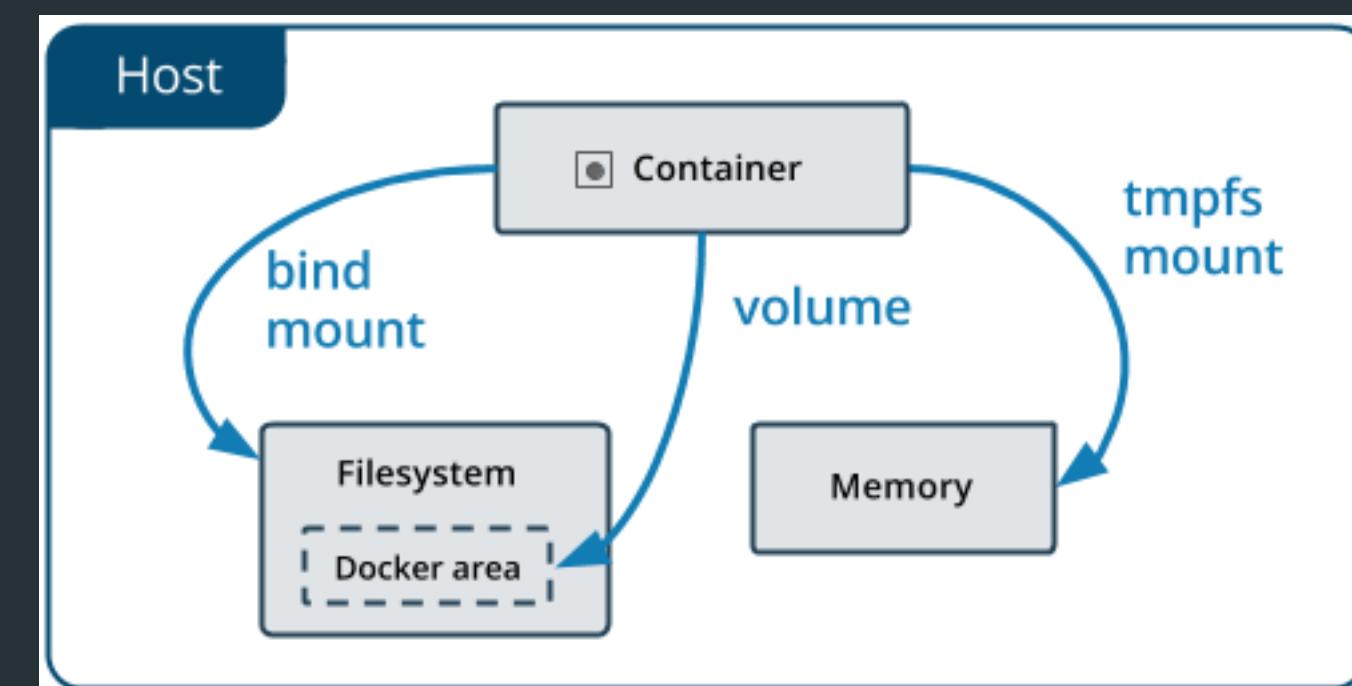
```
# Wait for a container to finish  
docker wait my-container
```

```
# Diff the changes between a container's image and its filesystem  
docker diff my-container
```

```
# Copy files from a container to the host  
docker cp my-container:/path/to/file /path/to/destination
```

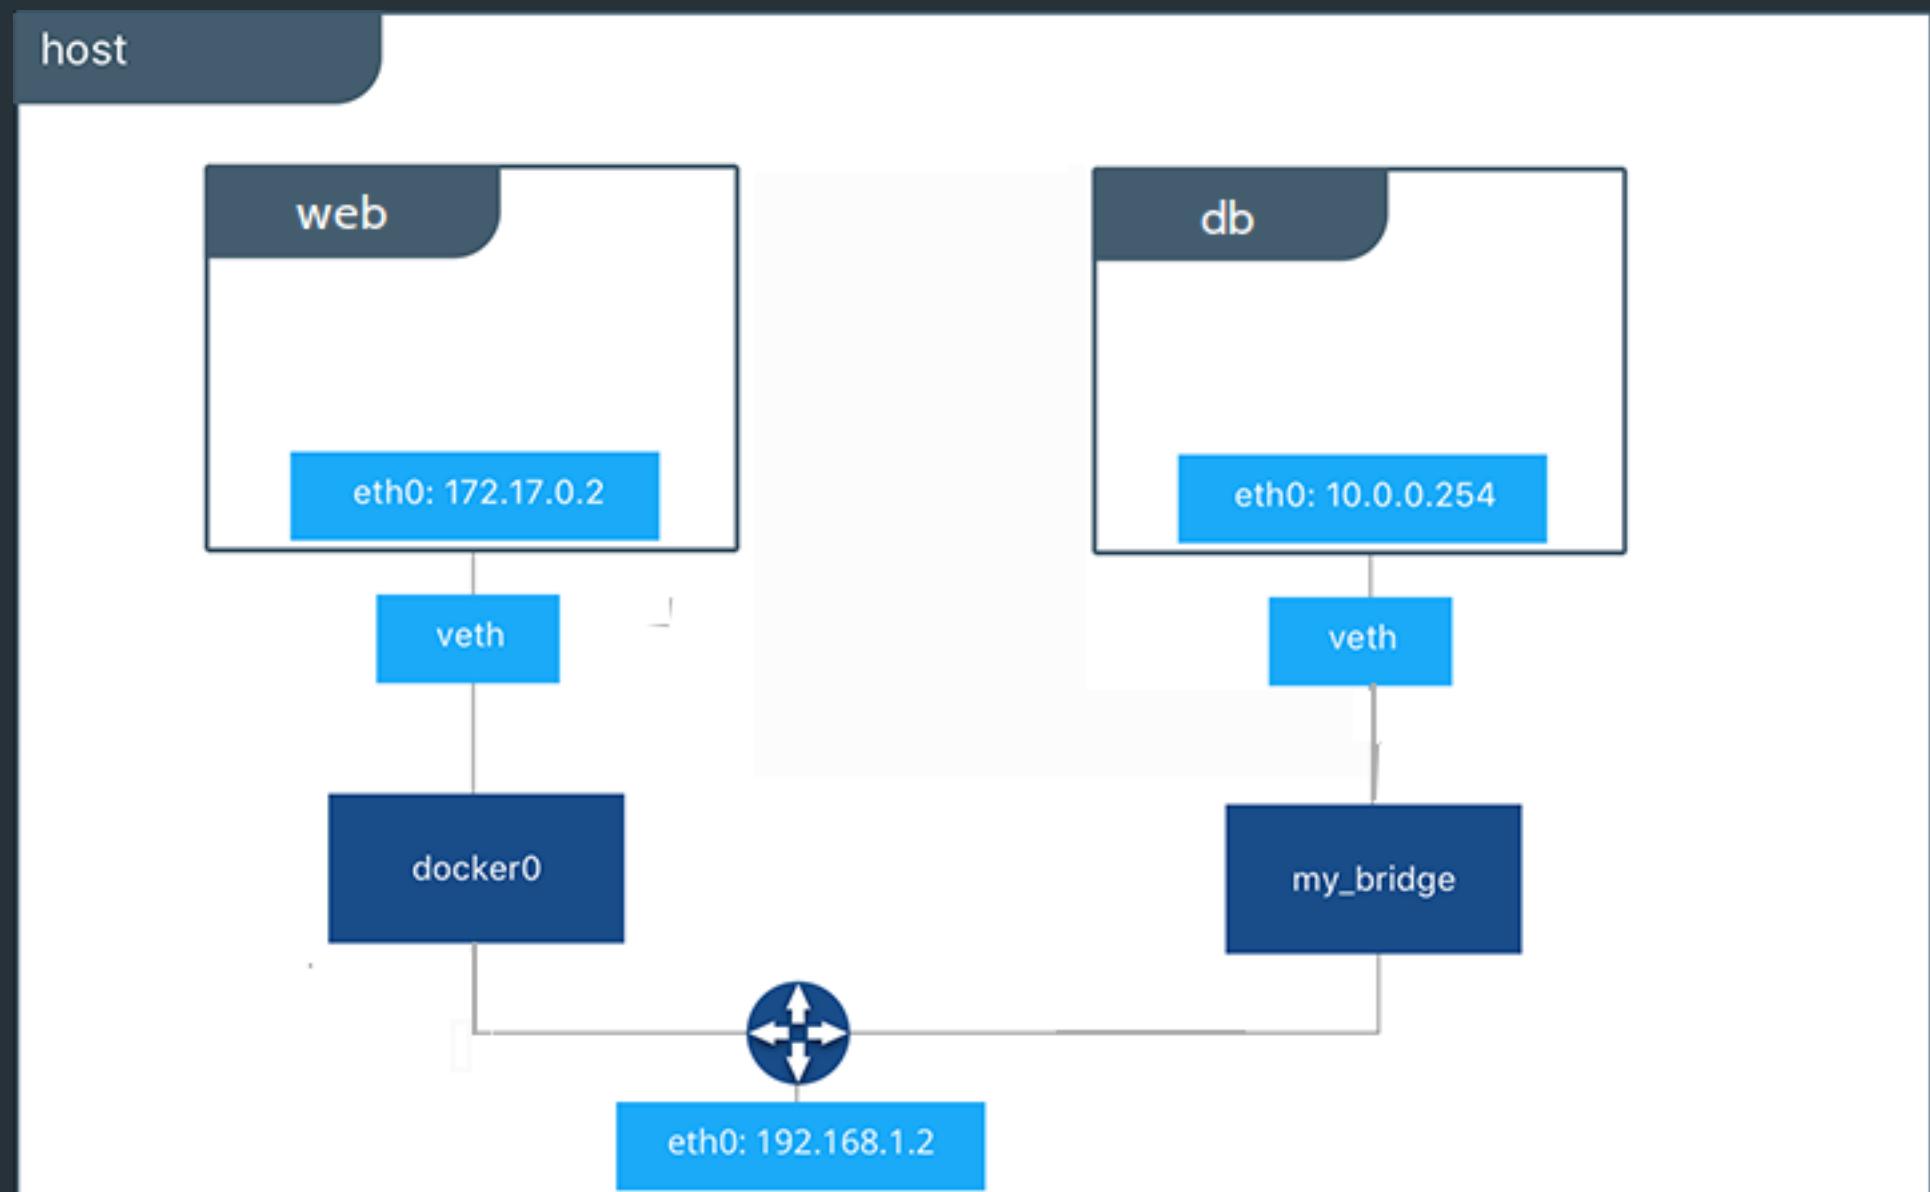
Connectivity and Storage

- Docker links: Docker links allow containers to communicate with each other by name.
- Docker networks: Docker networks allow containers to communicate with each other over a network.
- Docker volumes: Docker volumes allow containers to share data with each other and with the host machine.



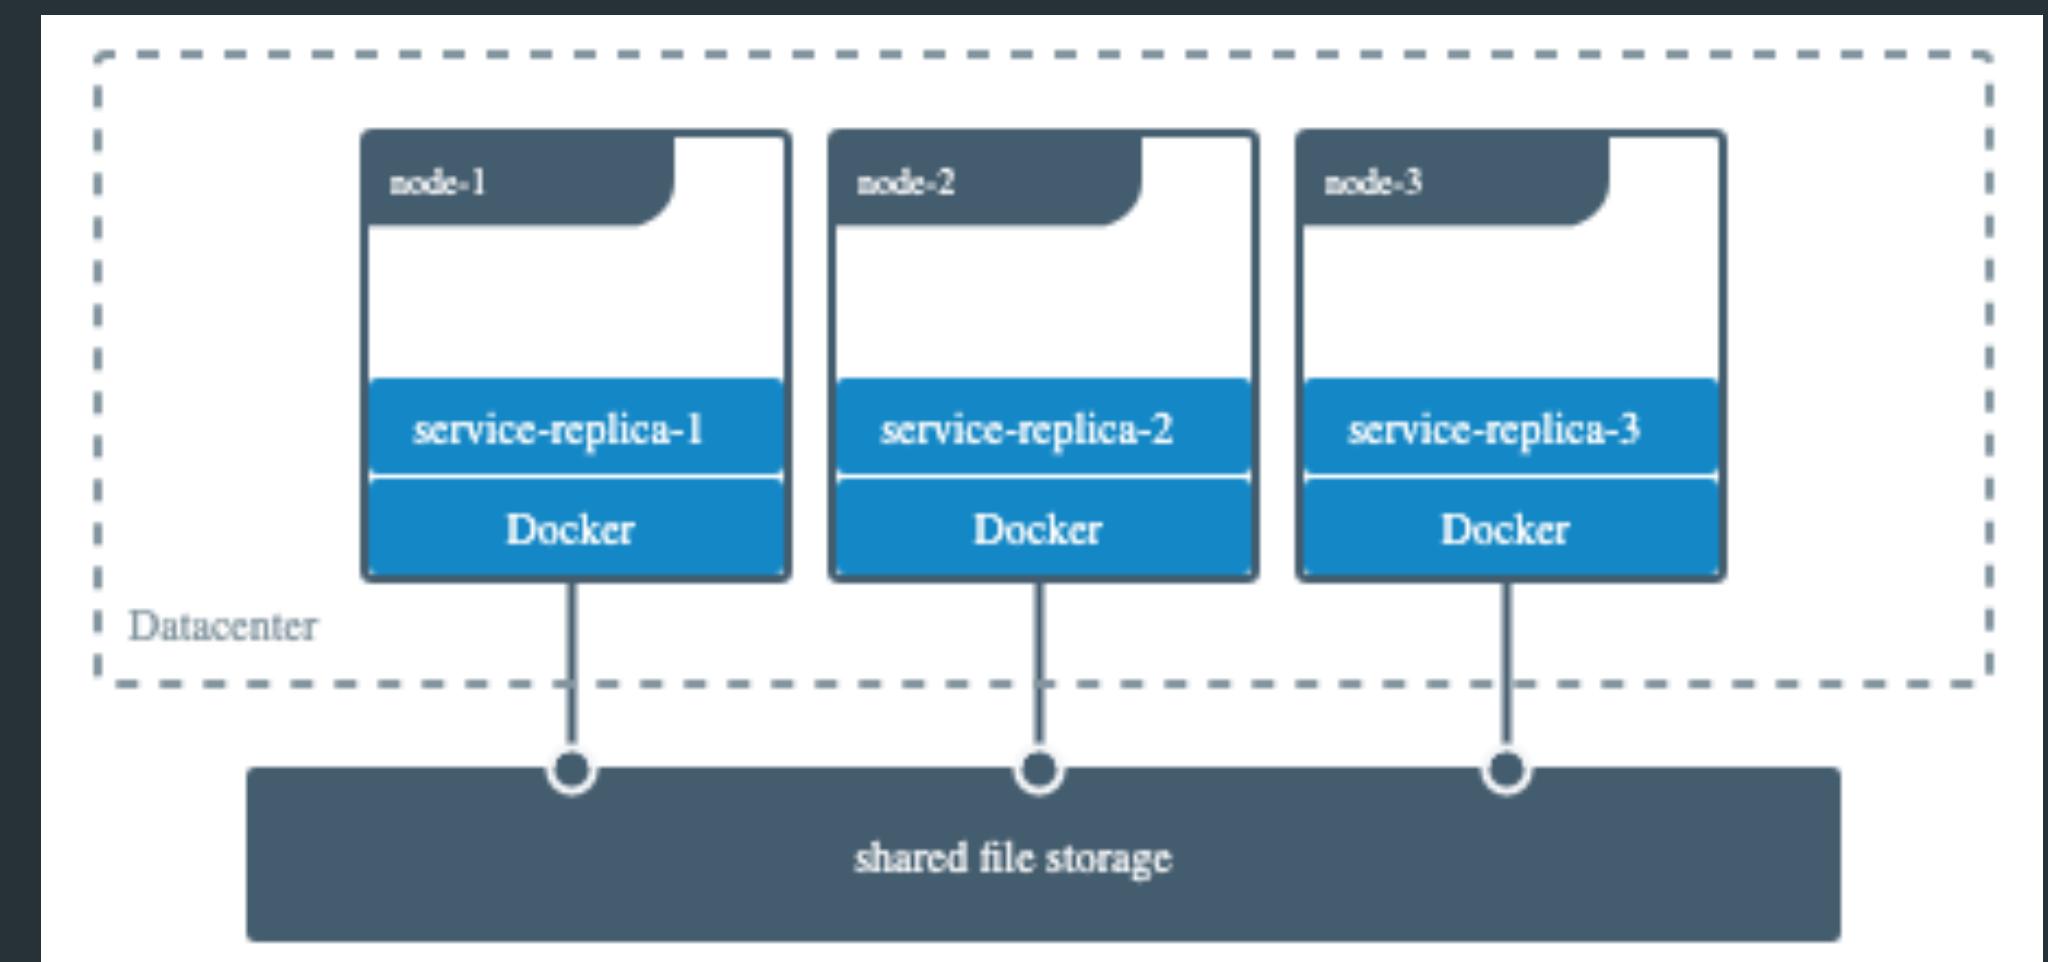
Docker Connectivity

- Docker containers can be connected to the host machine's network using a variety of methods, such as:
 - Bridged networking: Bridged networking allows containers to connect to the host machine's network as if they were physical machines.
 - Host networking: Host networking allows containers to share the host machine's network interface.
 - Overlay networks: Overlay networks allow containers to connect to a network that is shared by multiple hosts.



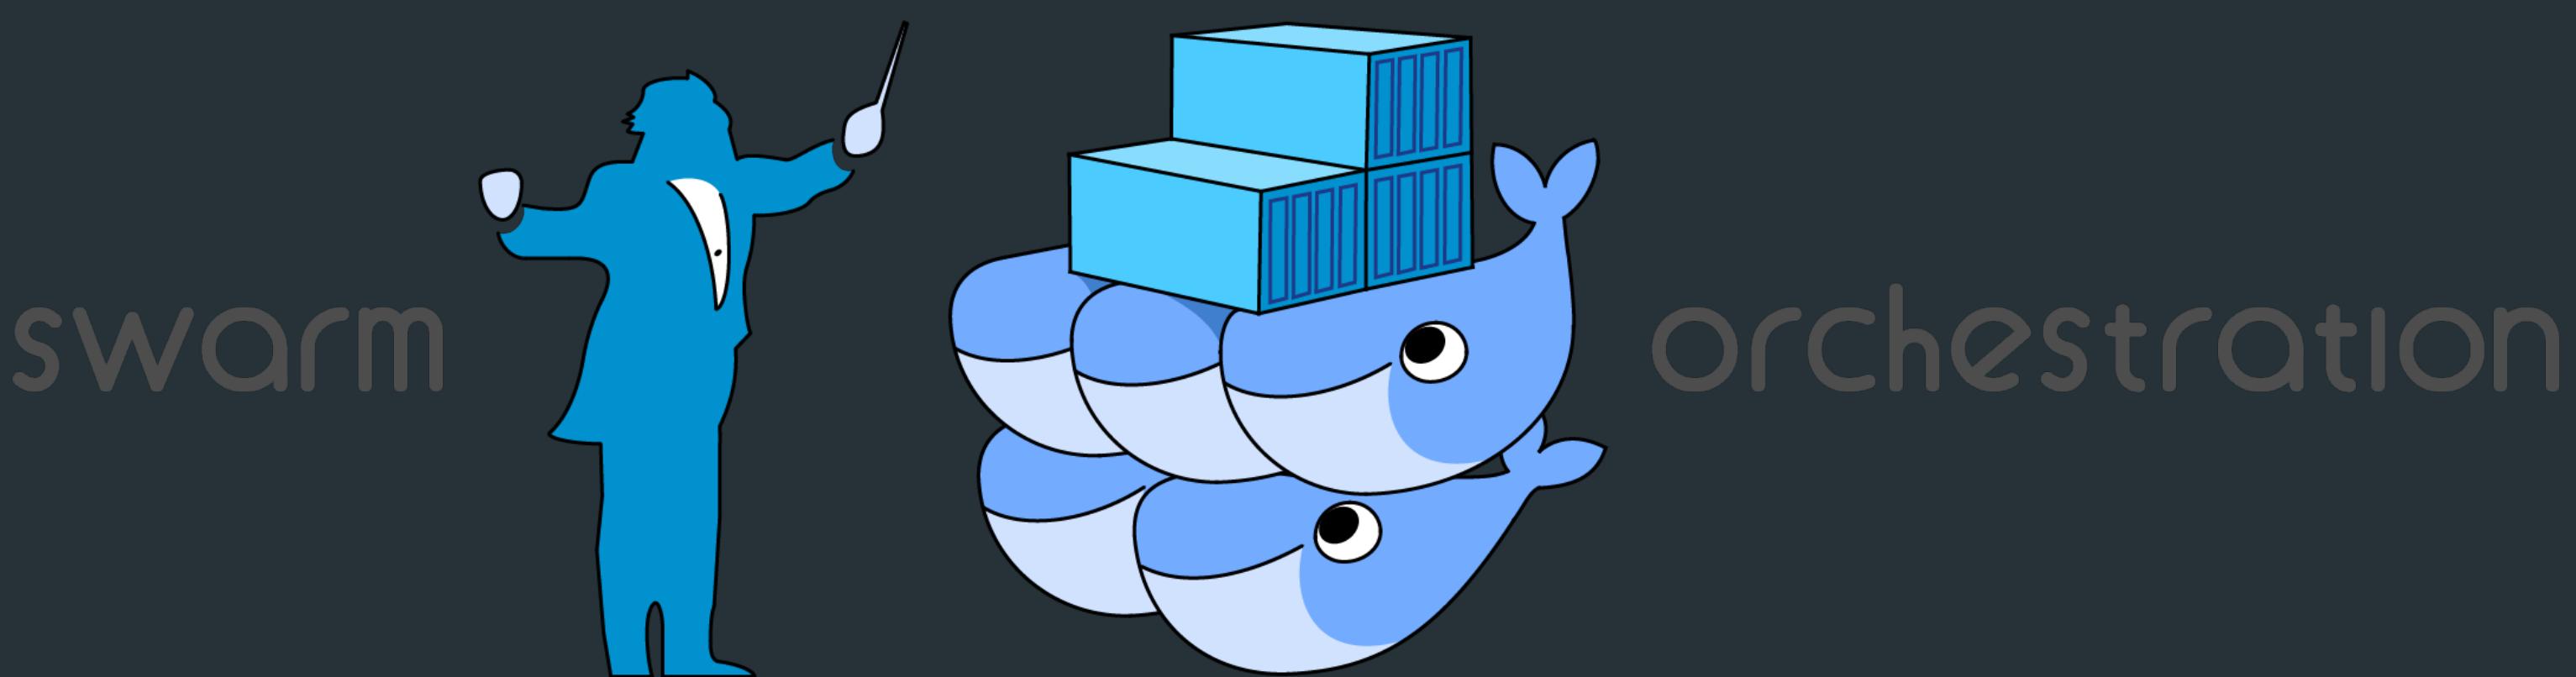
Docker Storage

- Docker volumes can be created as:
 - Local volumes: Local volumes are stored on the host machine's filesystem.
 - Named volumes: Named volumes are stored in a shared location, such as a Docker registry.
 - Tmpfs volumes: Tmpfs volumes are stored in temporary memory.



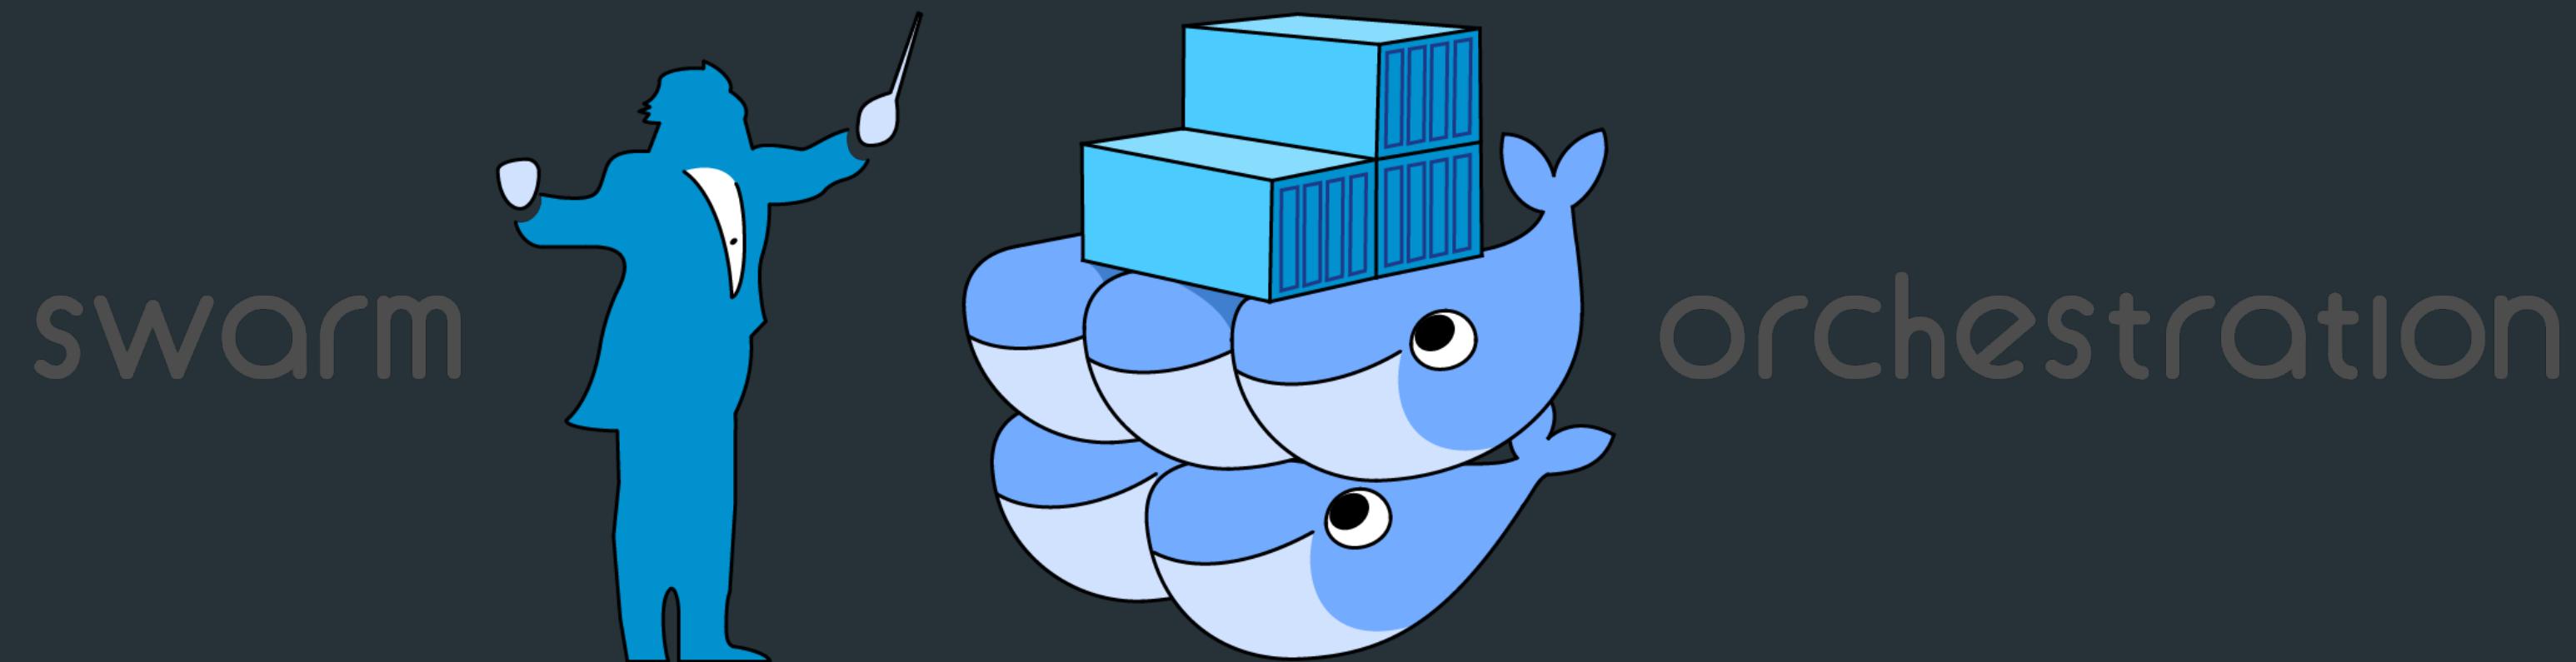
Docker Swarm

- A container orchestration system that allows you to run and manage Docker containers across multiple machines.
- A powerful tool that can be used to deploy and manage large-scale applications.
- Easy to use and can be managed using the Docker CLI or the Docker Swarm UI.



Benefits of Using Docker Swarm

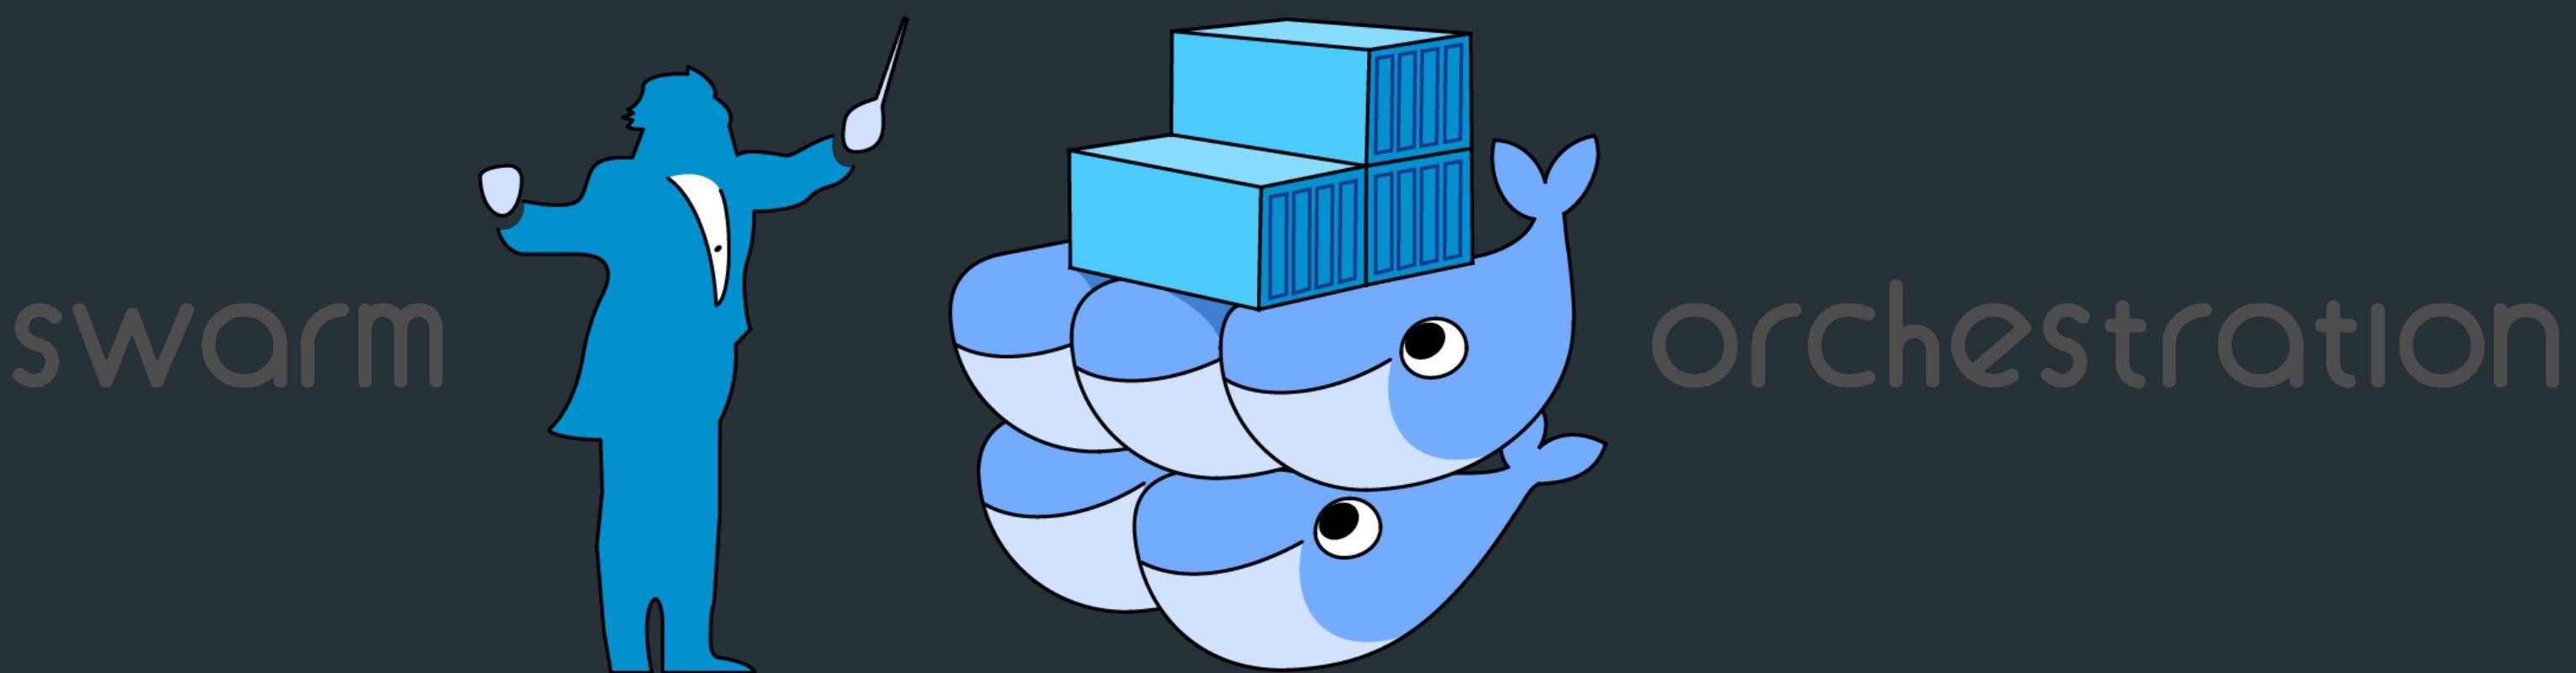
- Scalability
- High availability
- Ease of use
- Cost-effectiveness



How to Use Docker Swarm

- To use Docker Swarm, you first need to create a Docker Swarm cluster.
- Once you have created a Docker Swarm cluster, you can then deploy applications to the cluster.
- You can manage Docker Swarm clusters using the Docker CLI or the Docker Swarm UI.

`docker swarm init`

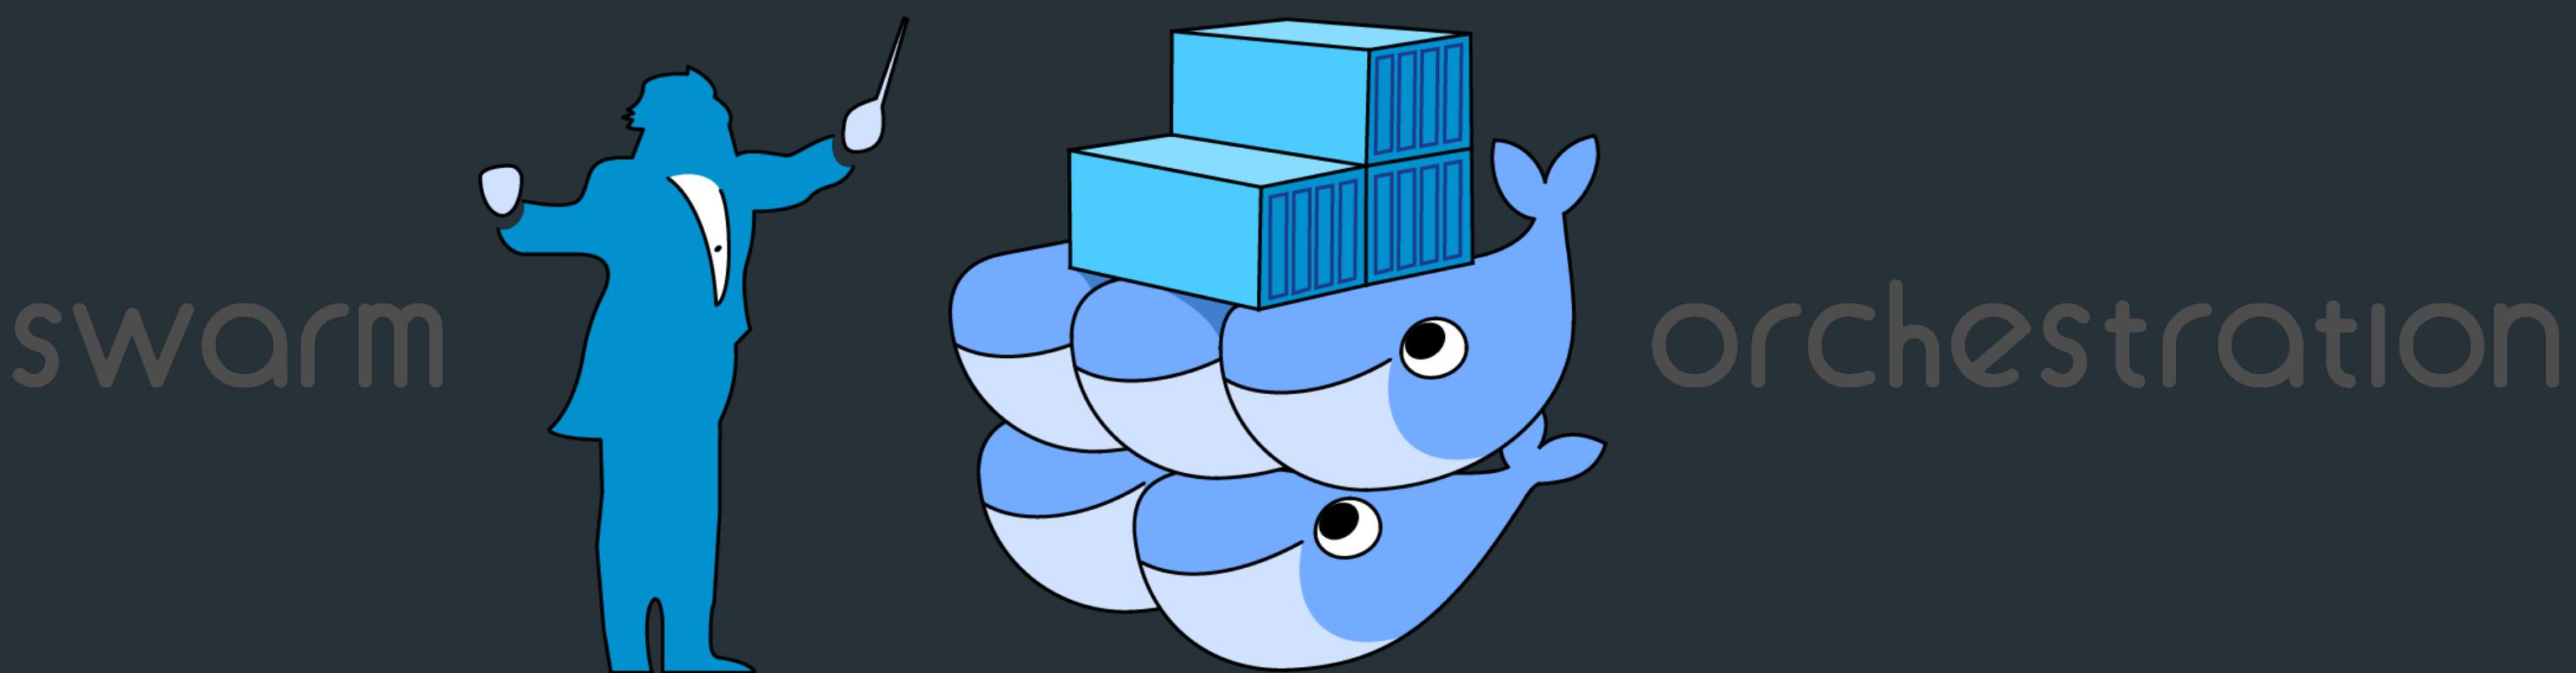


How to Use Docker Swarm

- To use Docker Swarm, you first need to create a Docker Swarm cluster.
- Once you have created a Docker Swarm cluster, you can then deploy applications to the cluster.
- You can manage Docker Swarm clusters using the Docker CLI or the Docker Swarm UI.

`docker swarm init`

`docker stack deploy -c <stack-file> <stack-name>`



Introduction to Kubernetes

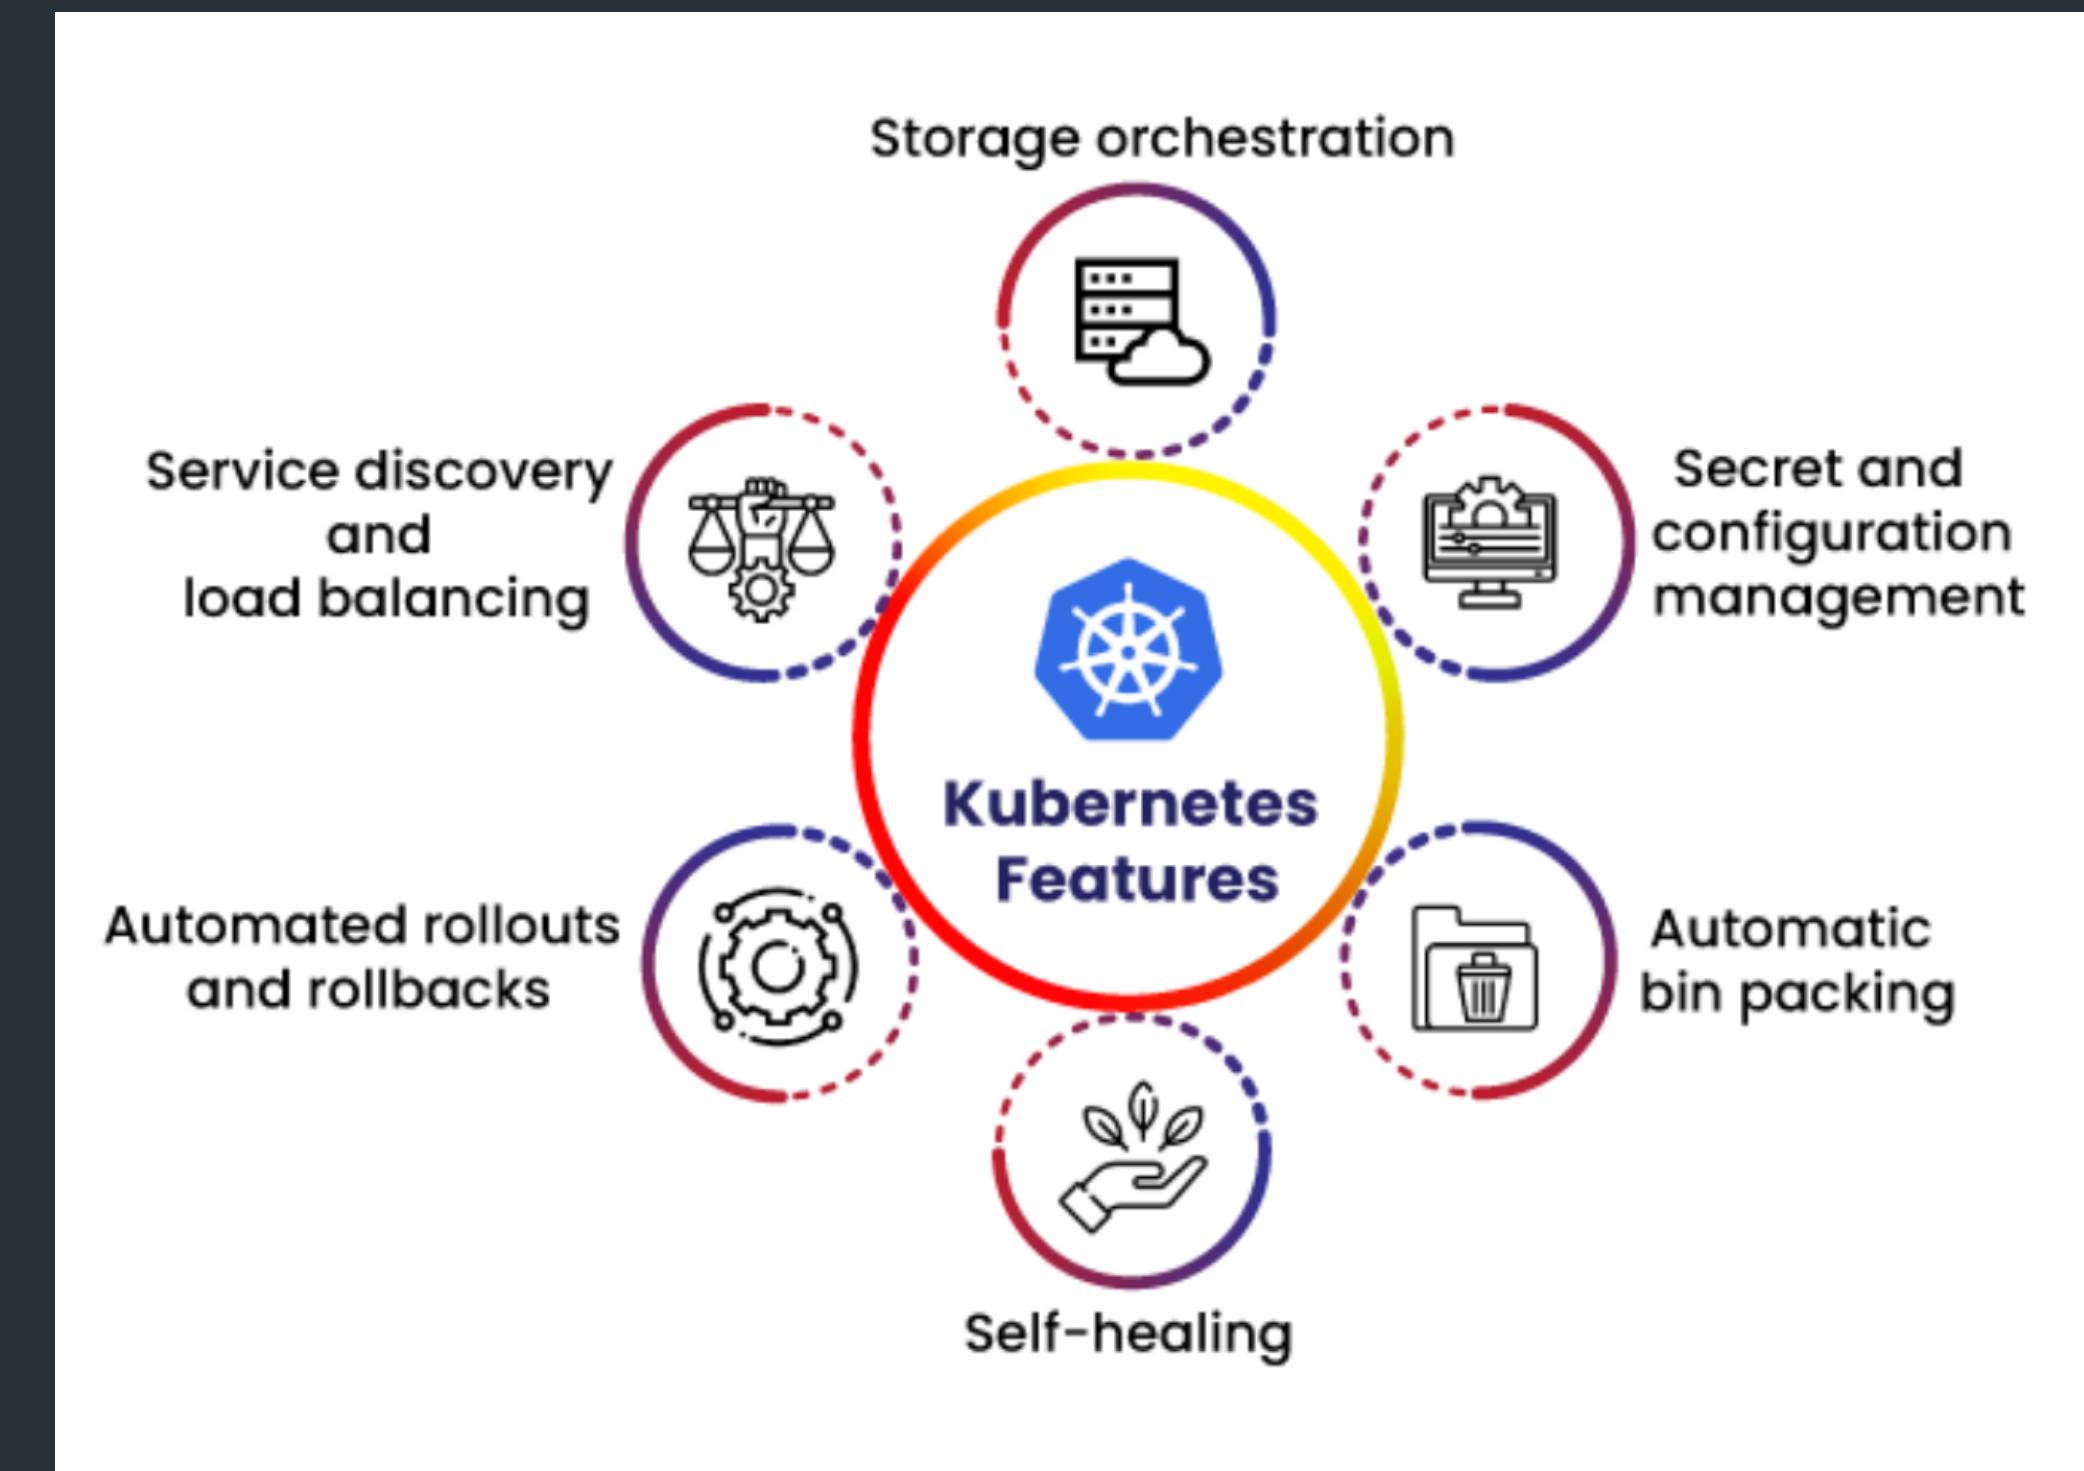
- An open-source container orchestration system.
- Used to automate deployment, scaling, and management of containerized applications.
- Used to deploy applications on a single machine or on a cluster of machines.



kubernetes

Benefits of Using Kubernetes

- Scalability: Kubernetes is highly scalable and can be used to deploy and manage applications of any size.
- High availability: Kubernetes is highly available and can continue to run applications even if some of the underlying nodes fail.
- Ease of use: Kubernetes is easy to use and can be managed using a variety of tools, including the command-line, the web UI, and the kubectl command-line tool.
- Cost-effectiveness: Kubernetes is cost-effective and can be used to save money on hardware and infrastructure costs.



How to Use Kubernetes

- To use Kubernetes, you first need to create a Kubernetes cluster.
- Once you have created a Kubernetes cluster, you can then deploy applications to the cluster.
- You can manage Kubernetes clusters using the command-line, the web UI, or the kubectl command-line tool.

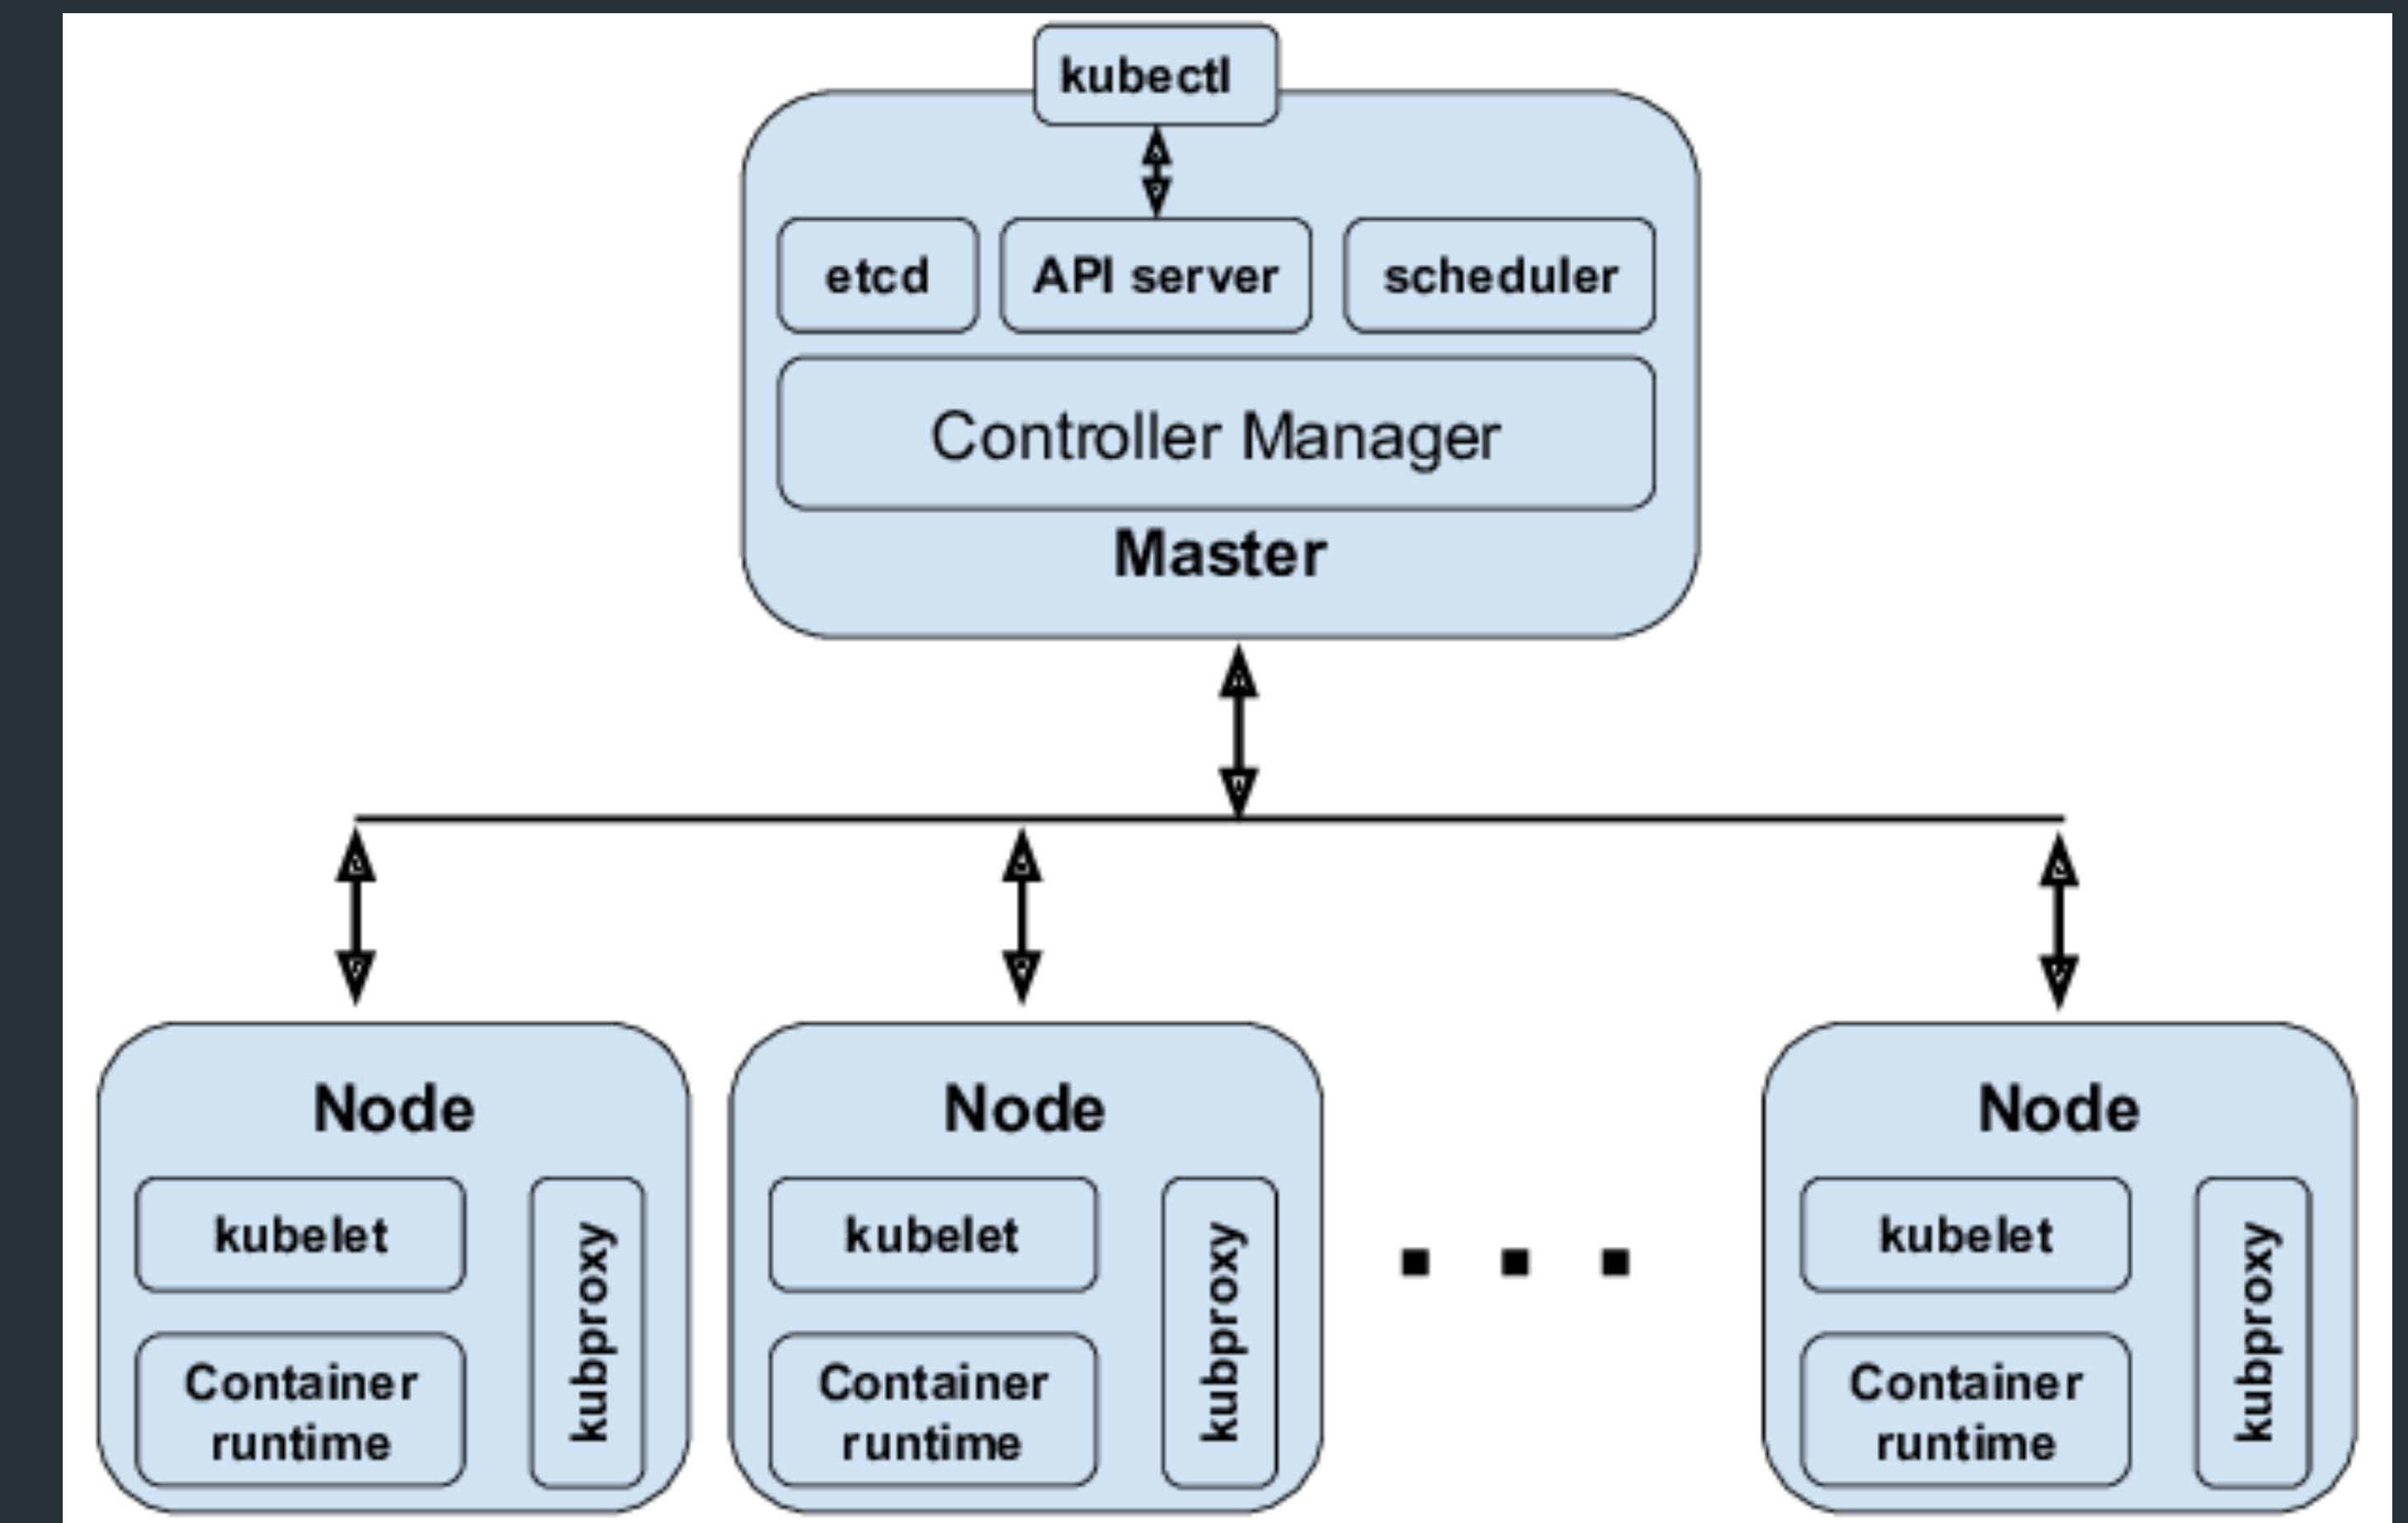
`kubectl create cluster`

`kubectl apply -f <manifest-file>`

Kubernetes Components

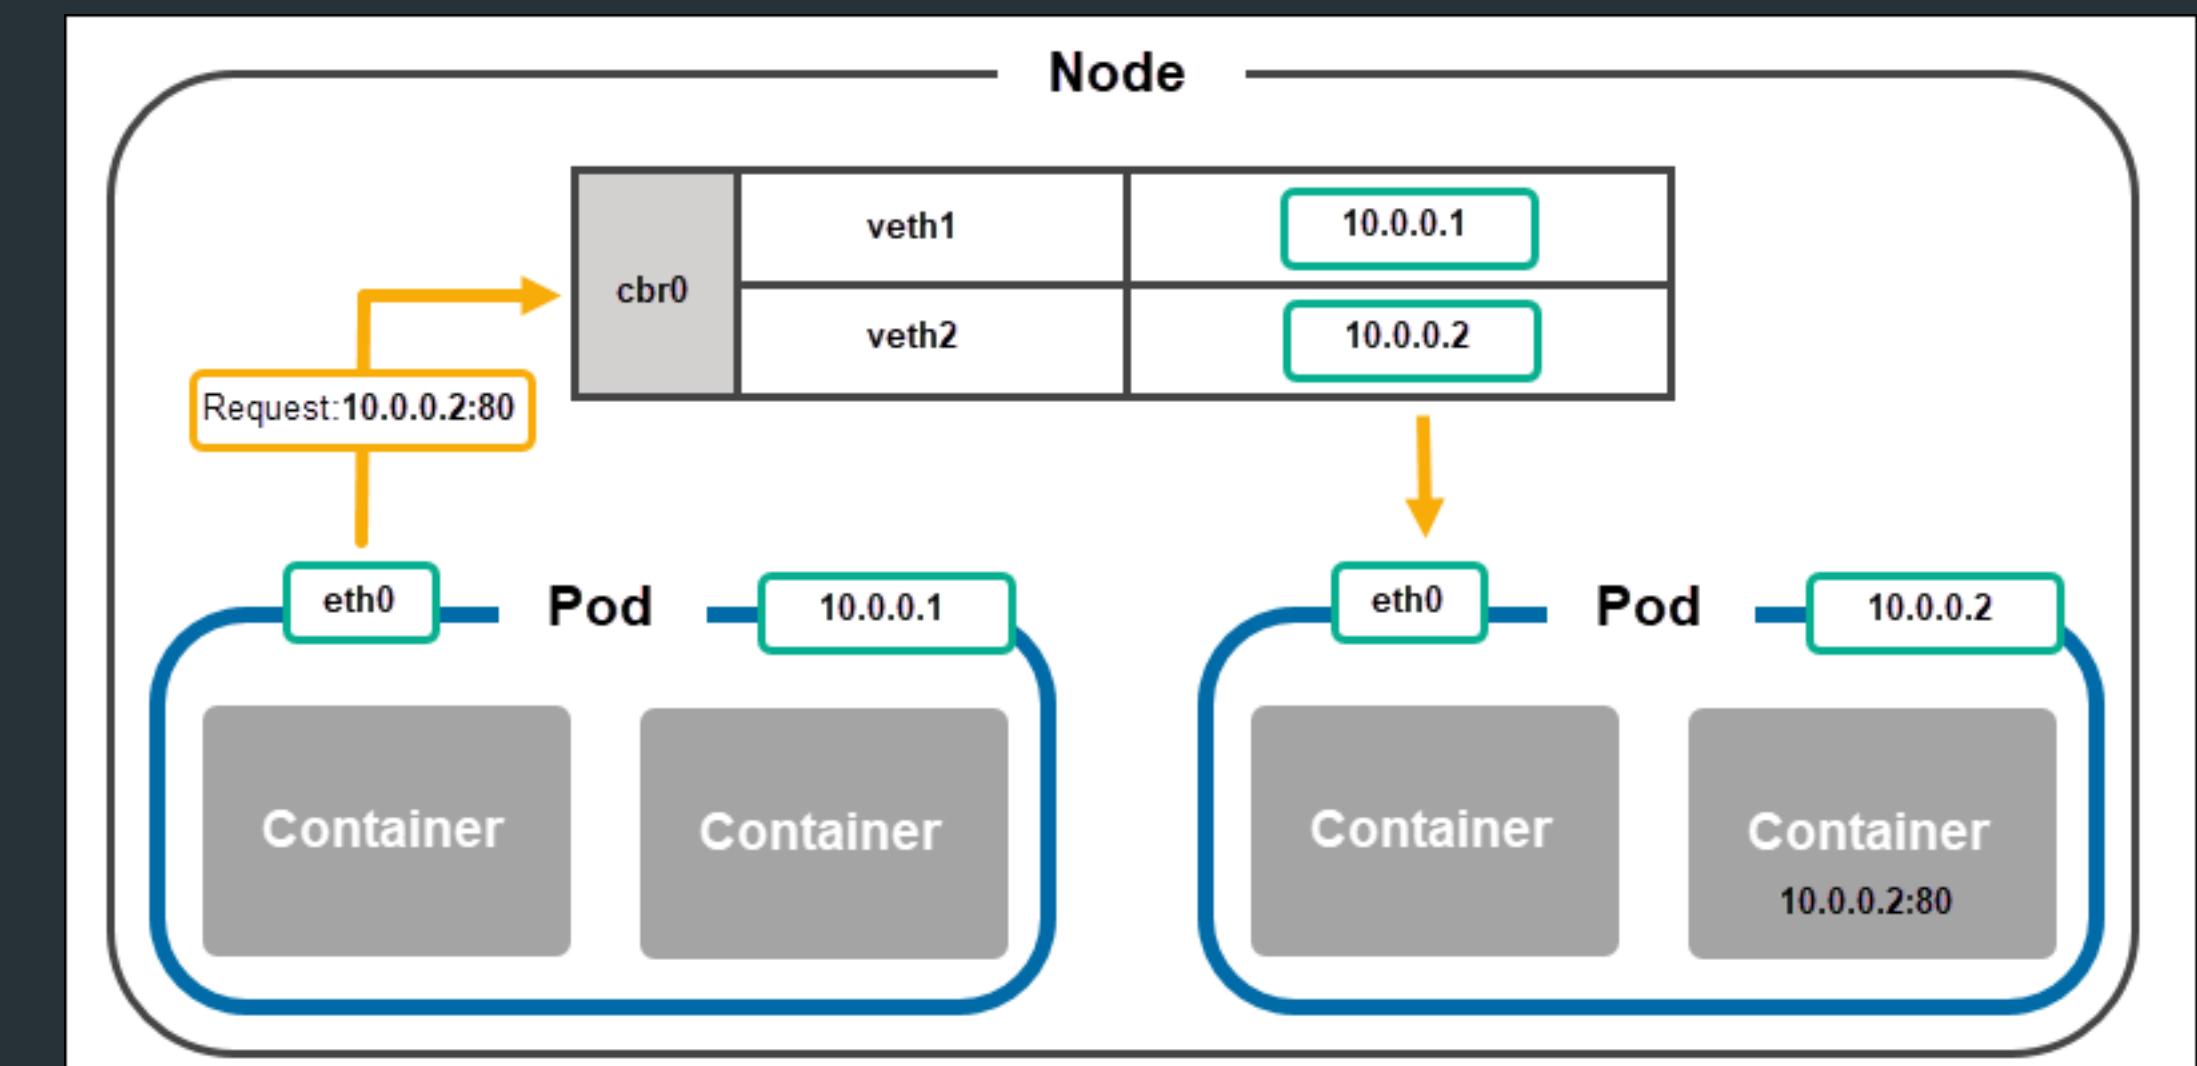
- The Kubernetes architecture is made up of the following components:

- Nodes
- Pods
- Services
- Controllers
- Schedulers
- API Server
- etcd
- Kubelet
- Kubectl



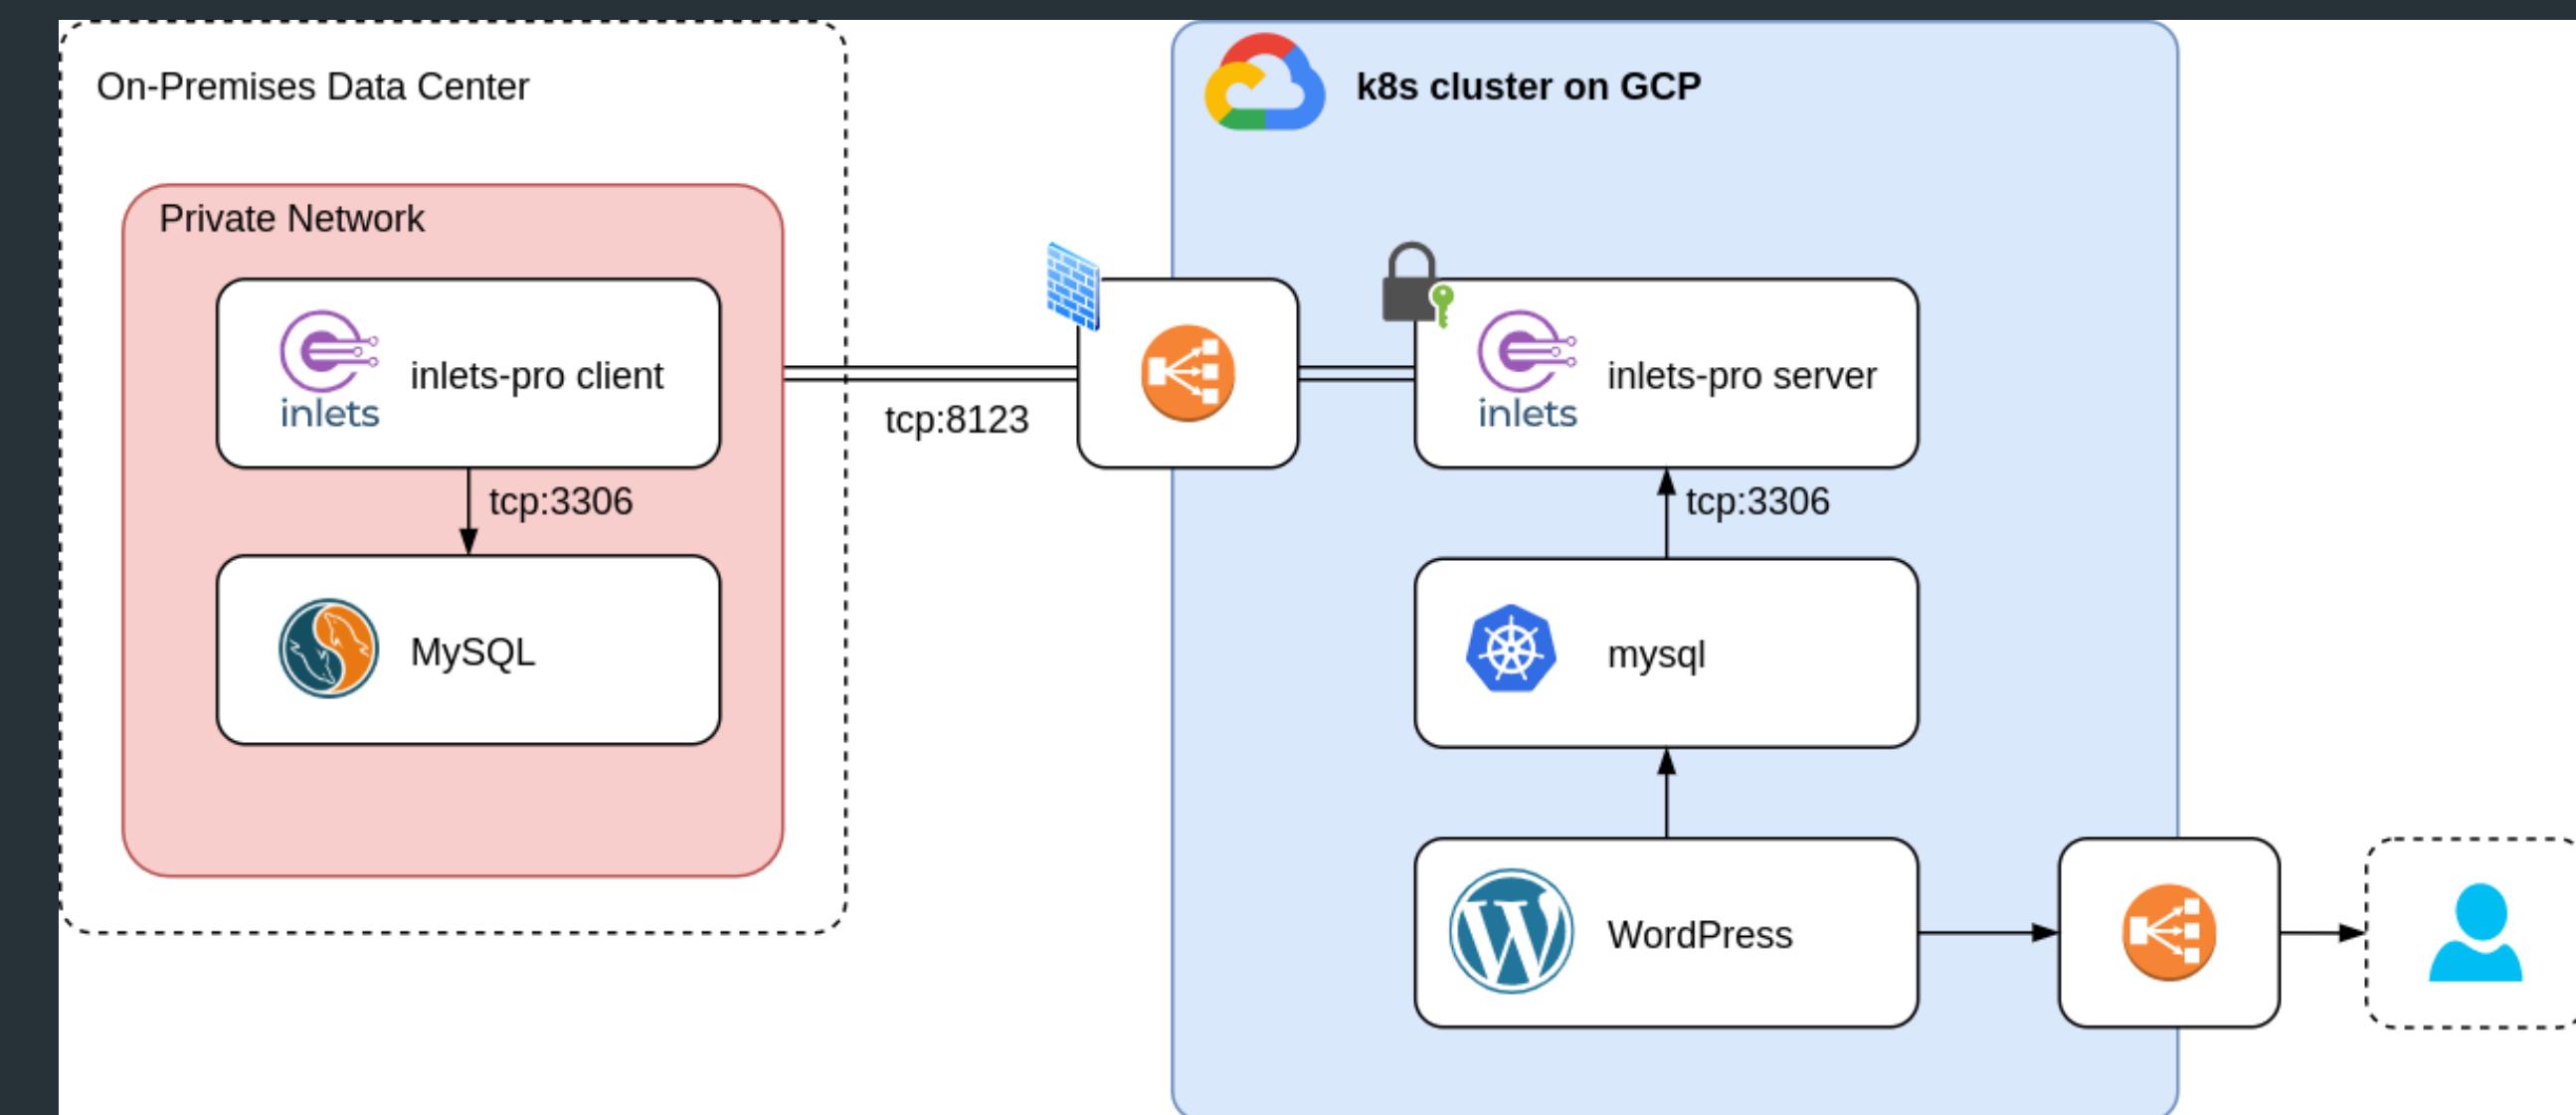
Kubernetes Communication

- Kubernetes components communicate with each other using a variety of protocols, including:
 - REST: The API Server uses the REST protocol to expose resources to clients.
 - gRPC: Controllers and kubelets use the gRPC protocol to communicate with each other.
 - etcd: etcd uses the Raft protocol to store configuration data.



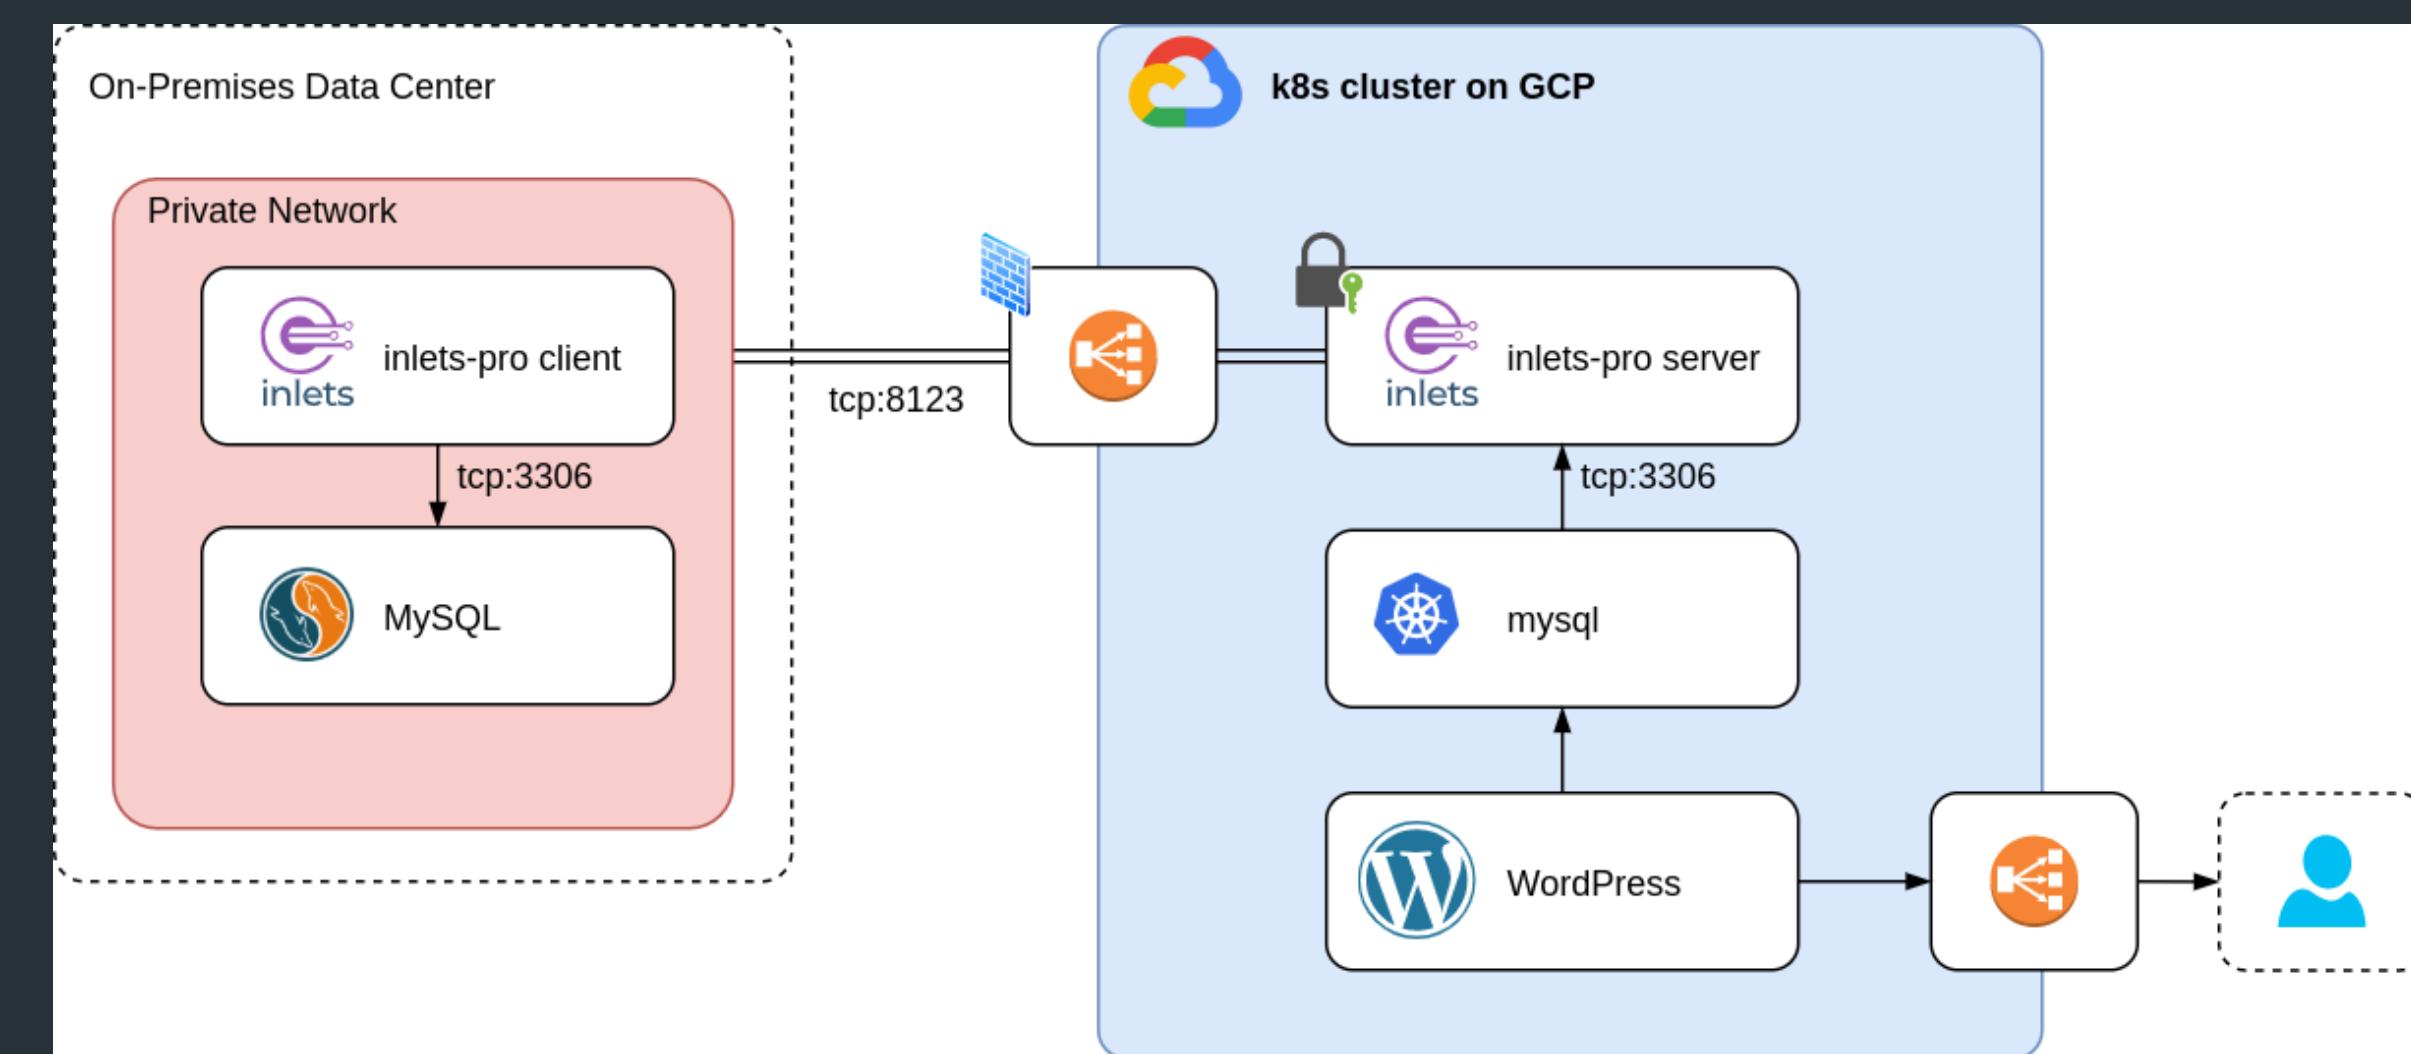
Setting Up a Kubernetes Cluster

- On-Premises
 - Are typically set up on a dedicated set of machines.
 - Offer the highest level of control and security.
 - Can be more expensive to set up and maintain than other types of clusters.



Setting Up a Kubernetes Cluster

- Cloud-Based
 - Are hosted by a cloud provider.
 - Typically more scalable and easier to manage than on-premises clusters.
 - Can be more expensive than on-premises clusters, depending on the provider and the features you need.



Setting Up a Kubernetes Cluster

- Managed
 - Provided by a third-party vendor.
 - Offer a turnkey solution for deploying and managing Kubernetes clusters.
 - Can be more expensive than self-managed clusters, but they can save you time and hassle.



Google Kubernetes Engine

Lecture outcomes

- Containers
- Virtualization
- Docker
- Kubernetes

